Definitional Instability in Kinase Inhibitor Selectivity: An Empirical Characterization and Motivated Properties for a Well-Formed Measure

Polina Vinogradova*,†,‡

†*Independent Researcher*

‡*ORCID: 0000-0003-3271-3841*

E-mail: polina.vino@gmail.com

Abstract

Kinase inhibitor selectivity is quantified using several competing definitions: S-score, selectivity entropy, Gini coefficient, and ratio-based measures, which are applied interchangeably without justification. Across four profiling datasets spanning three assay technologies (68-704 compounds each), we show that, with entropy and Gini computed on a common log-affinity scale, these definitions split three to one: the S-score, entropy and Gini agree closely while the fold-selectivity ratio agrees with none. Rank instability is associated with binding profile shape. We trace instability to three sources: convention-dependent scores for compounds with no binding above the activity threshold, near-tied top targets, and distribution-ratio disagreement for promiscuous compounds with a dominant target. We motivate four properties that a well-formed selectivity measure should satisfy. None of the four assessed satisfies all of them; the one property all four violate is a reliability condition for compounds without binding above threshold. We then propose a measure that does, recovering its full-panel ranking from about a third to a half of a panel, depending on the dataset.

Keywords: binding selectivity, cancer drug discovery, cheminformatics, kinase inhibition, research methodology.

Introduction

Kinase inhibitors are among the most intensely investigated compound classes in drug discovery, with 85 FDA-approved agents and several hundred in clinical development.[1] Because the human kinome is comprised of more than 500 proteins sharing a highly conserved ATP-binding site,[2] the design of inhibitors with controlled target profiles remains a central challenge. Selectivity is a primary criterion in lead optimization and a key determinant in the drug approval process.[3]

Despite its importance, no single consistent definition of selectivity exists. The field employs at least four distinct quantitative measures: the S-score,[4] selectivity entropy,[5] Gini coefficient,[6] and ratio-based measures, a family to which the field assigns the CATDS score.[7] These definitions are applied interchangeably in the literature without systematic justification, and their parameter values (the activity threshold in the S-score, the baseline in the entropy and Gini coefficient, the reference off-target in ratio-based measures) are chosen by convention rather than principle. Standardizing them is a prerequisite for integrating selectivity data across laboratories and platforms.

Bosc et al.[8] showed that these definitions can produce substantially different compound rankings on the same dataset, and the KInhibition platform[9] explicitly acknowledged that conflicting results across metrics are common. However, the structural sources of this disagreement, such as "which binding profile properties determine when definitions agree or conflict, and why?", have not been characterized.

The absence of a principled basis for selectivity measurement is not merely a practical inconvenience. Go/no-go decisions in lead optimization are made on the basis of selectivity assessments whose sensitivity to definitional choice is unknown. The lack of consistency

across applications of selectivity definitions for different compounds makes comparison very challenging.

Our analyses suggest that existing selectivity definitions are, in some cases, measuring entirely different qualities of binding profiles, and therefore cannot be easily converted from one to another. In information theory, Shannon[10] settled which function should measure entropy by writing down a few properties a sensible measure must have (e.g., that it should grow as the number of equally likely outcomes grows) and showing that only one formula satisfies them. In this work, we follow this precedent, characterizing how the competing measures actually behave, stating the properties a good measure should have, and proposing a satisfactory formula.

We use four large profiling datasets for kinase inhibitor selectivity that span three assay technologies in our characterization. In service of the overarching goal of developing a principled, standardized formula for defining selectivity, our contributions are as follows:

(i) we show that the S-score, entropy and Gini agree closely with each other on every one of the four datasets (Spearman 0.72–1.00), while the fold-selectivity ratio $R(k)$ agrees with none of them (0.14–0.62). We call the first group the distribution-based family, and $R(k)$ the ratio-based family. The split follows what each quantifies: how concentrated a compound's activity is across the whole panel, against how large the gap is to one chosen off-target. This grouping holds for entropy and Gini computed on a common baseline-subtracted log-affinity scale; under the association-constant form of the original publication, the entropy does not cluster with the S-score and Gini on any of the three affinity datasets (Methods);

(ii) we identify three sources of instability: scores set by an arbitrary convention for compounds with no binding above the activity threshold, ratio-specific instability driven by near-tied top targets, and distribution-ratio disagreement for compounds with one dominant target and broad moderate off-target activity;

(iii) we propose four motivated properties: (D1) a reliability threshold condition, (D2) bounded gap sensitivity, (D3) distributional consistency, and (D4) monotonicity under weak off-target addition. We then find that none of the four definitions assessed satisfies all four simultaneously. The one property all four violate is D1, a reporting convention;

(iv) we propose a candidate measure, a gated, smooth-hinge effective-target number, to show that the four properties can be met together. It recovers its full-panel ranking from roughly a third to a half of a typical panel, depending on the dataset, and we prove that, unlike the entropy, it cannot be made to look more selective by adding a weakly bound kinase.

We note that the candidate measure we propose is neither unique as a measure satisfying the four properties jointly, nor is it guaranteed to perform better than existing measures in all cases.

Background

Existing Selectivity Metrics and Their Comparison

Several metrics have been proposed to quantify kinase inhibitor selectivity from profiling panel data. Karaman et al.[4] introduced the selectivity score (S-score), defined as the fraction of tested kinases bound with a dissociation constant below a fixed cutoff. Uitdehaag and Zaman[5] proposed the selectivity entropy, which treats the normalized binding profile as a probability distribution and computes its Shannon entropy. Low entropy indicates activity concentrated on few targets. Graczyk[6] adapted the Gini coefficient from economics, treating the distribution of a compound's activity across the panel as an inequality distribution. A value near 1 indicates activity concentrated on few kinases and near 0 indicates activity spread evenly. Klaeger et al.[7] introduced the concentration- and target-dependent selectivity (CATDS) score, which quantifies the reduction in binding to a given target relative to the

summed reduction across all detected targets at a specified inhibitor concentration (see Discussion for why we did not implement it as a fifth measure).

Bosc et al.[8] subsequently proposed two further metrics, the window score and ranking score, and compared them against four previously published metrics (the S-score, Gini coefficient, selectivity entropy, and a partition index) on the Davis and Anastassiadis datasets, showing that compound rankings differ substantially depending on the metric applied and that certain compounds cannot be scored under some definitions when threshold conditions are not satisfied. For example, ratio-based scores are undefined for a compound with no kinase above the activity threshold, since there is then no primary target or off-target to compare. Our own implementation of $R(k)$ does not follow this convention (see Methods, "Ratio"), which bears on $R(k)$'s D1 grading in Table 3.

Reconciliation attempts have remained limited in scope. Bajorath et al.[11] compared CATDS-based selectivity assessments from Klaeger et al. with activity data mined from ChEMBL for the same compounds, finding broad consistency between the two sources. The KInhibition platform[9] explicitly acknowledged that multiple existing metrics yield conflicting results for the same compound, but addressed this by introducing its own composite score (combining on-target inhibition with an off-target inhibition penalty) rather than reconciling the disagreement among the existing definitions. More recently, the same group built a KIRHub resource[12] on a new profiling campaign of 86 approved kinase inhibitors against 758 kinases and disease-associated mutant variants. This resource offers users a choice between a further named metric (the KISS score, an on-target/off-target balance) and a separate Gini-based score, thus proliferating named, ad hoc definitions.

Two further notable gaps surface in existing work. The first is that each metric is evaluated at fixed parameter values. Thus, the sensitivity of rankings to parameter choice has not been systematically characterized. The second is that the structural properties of binding profiles that determine when definitions agree or disagree have not been identified. The present work addresses both gaps.

Axiomatic Approaches to Ranking and Measurement

Precedent for an axiomatic characterization of measurement functions is established. Shannon[10] demonstrated that the entropy of a probability distribution is uniquely determined, up to a multiplicative constant, by three requirements: continuity in the probabilities, monotonic growth with the number of equally likely outcomes, and consistency when outcomes are grouped. Khinchin[13] subsequently gave an alternative axiomatization (continuity, maximality on the uniform distribution, expansibility, and additivity over independent subsystems). The selectivity entropy of Uitdehaag and Zaman applies the Shannon formula to normalized binding profiles but provides no analogous axiomatic justification. The axioms that single out Shannon entropy, such as maximality on the uniform distribution and additivity over independent subsystems, characterize a measure of uncertainty rather than a measure of selectivity.

Arrow[14] established that no social welfare function satisfies simultaneously the requirements of unrestricted domain, Pareto efficiency, independence of irrelevant alternatives, and non-dictatorship. The theorem is relevant here in two respects: it demonstrates that axiomatic analysis of a ranking concept can yield impossibility, and it shows that competing ranking functions reflect genuine tradeoffs between properties that cannot all be achieved simultaneously. Singleton and Booth[15] applied this methodology to truth discovery, i.e. the problem of identifying true facts from conflicting reports of unknown reliability, posing it as a social choice problem, formulating candidate axioms, and proving that a subset cannot be jointly satisfied.

To our knowledge, this axiomatic style of analysis has not previously been applied to any selectivity or promiscuity measure in drug discovery. The present work proposes empirically motivated properties as a foundation for such an analysis.

Methods

Datasets

Our primary analysis uses two publicly available kinase inhibitor profiling datasets, described first. Two further datasets are added below to confirm the findings across additional assay technologies. The Davis dataset[16] is comprised of binding affinities (K_d) for 68 small-molecule kinase inhibitors measured against 433 human kinase domains. Affinity values are reported as dissociation constants and were converted to $\mathrm{p}K_d = -\log_{10}(K_d/10^9)$ prior to analysis, following the convention of the DeepDTA benchmark.[17] The resulting affinity matrix is complete (68×433, 29,444 entries), with all non-binding pairs assigned the detection-floor value $\mathrm{p}K_d = 5.0$ (corresponding to a K_d of 10 μM). This is the standard convention for the dataset[8,17] and slightly underestimates true non-binding affinity.

Active entries (pKd > 6, i.e. $K_d < 1$ μM) make up 17.9% of the Davis matrix (`dataset_descriptive_stat`
This cutoff answers "was binding detected above assay noise/floor," the conventional boundary used in kinase profiling, rather than "is this target plausibly occupied at physiological drug concentrations," a question affinity values alone cannot answer without exposure data (e.g. free-drug concentration) that neither dataset provides. A stricter cutoff ($\mathrm{p}K_d \geq 7$, $K_d < 100$ nM) shrinks per-compound active sets substantially (mean active kinases per compound falls from 77.5 to 34.4, a 2.25-fold reduction) and pushes two of the 68 Davis compounds, which have at least one active kinase at $\mathrm{p}K_d > 6$, to zero active kinases at the stricter cutoff (into the zero-active regime discussed in Results). Shifting the S-score's threshold sweep up by one log unit to match this cutoff (from $[5.5, 8.0]$ to $[6.5, 9.0]$; the entropy, Gini, and ratio sweeps are unchanged by this check) lowers the S-score's correlation with entropy and Gini from ≈ 0.95 to ≈ 0.83 and raises its correlation with the ratio from 0.456 to 0.574. The S-score therefore stays much closer to entropy and Gini than to the ratio (`dataset_descriptive_stats.py`).

The Klaeger dataset[7] is comprised of binding affinities for 243 clinical kinase inhibitors

profiled against the human kinome using Kinobeads chemical proteomics. Apparent dissociation constants (K_d^{app}) were derived from dose–response curves across eight inhibitor concentrations in cancer cell lysates. Only high-confidence measurements (as annotated by the original authors) were retained, yielding 5,123 drug–target pairs for 222 drugs and 343 kinase targets. Affinity values were converted to pK_d using the same formula as above. Missing entries, representing kinase–inhibitor combinations not detected in the assay, were assigned the background value p$K_d = 5.0$. The resulting Klaeger matrix is 222×343. Active entries (pKd > 6) make up 3.6% of the Klaeger matrix, versus 17.9% for Davis above (`dataset_descriptive_stats.py`). The lower active fraction for Klaeger reflects the more stringent detection threshold of the chemoproteomic assay.

The high-confidence filter removes 793 of the 5,916 drug–target pairs with any reported K_d^{app} (13.4%). A removed pair is treated identically to an untested one and assigned the same p$K_d = 5.0$ background value. Relaxing the filter (retaining the 793 low-confidence pairs alongside the 5,123 high-confidence ones) changes only 0.74% of the resulting matrix's entries, shifts every cell of Table 1's Klaeger triangle by at most 0.05, and leaves every panel-size p^* value reported in Figure 5 and Figure 7 unchanged. The filter choice is therefore not responsible for the reported correlations or convergence thresholds.

To test whether our findings depend on the assay technology, we added two further datasets that measure kinase inhibition by different means. The Anastassiadis dataset[18] profiles 178 inhibitors against 300 kinases in a single-concentration (0.5 μM) radiometric activity assay, reported as percent residual kinase activity relative to a solvent control. We converted these to percent inhibition (100 minus residual activity, clipped to $[0, 100]$, where higher values denote stronger inhibition), with the 1.1% missing entries set to zero inhibition (`dataset_descriptive_stats.py`). The Metz dataset[19] reports pK_i affinities for a large medicinal-chemistry campaign using the same negative-log scale as pK_d. The value p$K_i = 4.0$ is the lowest value reported anywhere in the released data (only 0.17% of measurements), so we assign it to unmeasured combinations as the assay's own detection floor

(the same convention used for the Davis $pK_d = 5.0$ floor above). Coverage per compound is highly skewed (a median compound is measured against only 21 of the 172 kinases; `dataset_descriptive_stats.py`). Because of this, we restricted the dataset to the 704 compounds profiled against at least 50 of the 172 kinases ($\geq 29\%$ coverage). The threshold excludes the most sparsely profiled compounds while retaining a large enough compound set for the panel-size and instability analyses. It does not make most entries real measurements: across the 704 retained compounds, coverage ranges from 29.1% to 83.1% of the 172 kinases (mean 40.9%, median 37.2%, i.e. 64 of 172 kinases), so for the median retained compound most entries are floor-imputed, and fewer than a quarter of compounds have more than half of their entries measured (75th percentile 47.1%). A sizable minority sit close to the threshold: 23 of 704 (3.3%) have exactly 50 measured kinases, and 170 (24.1%) have 55 or fewer (`metz_coverage.py`).

The four datasets span three assay technologies (competition binding, chemoproteomics, and functional activity) and two readout types (affinity and single-concentration inhibition). Each selectivity definition is computed on the native readout of its dataset, and our analyses concern rank correlations between definitions and rank instability. Figure 1 shows the definitions are not invariant to how they are operationalized on a given readout: two defensible operationalizations of the entropy rank compounds differently, and on Klaeger the relationship between them changes sign ($r = -0.32$, a weak negative association, not a near-complete inversion of the ranking). Cross-dataset comparison here is based on rank correlations computed within each dataset. For the percent-inhibition and pK_i datasets, the S-score threshold, and entropy/Gini baseline grids were shifted to span each dataset's active range (entropy/Gini baselines over $[0, 30]\%$ inhibition on Anastassiadis and $pK_i \in [4.0, 5.5]$ on Metz, against $[5.0, 6.5]$ on the pK_d datasets), while the ratio off-target index $k \in \{1, \ldots, 5\}$ was left unchanged (it is an index of off-targets and so needs no rescaling).

We excluded two additional candidate datasets. The Gao screen[20] is a single-concentration functional inhibition assay of the same readout type as the Anastassiadis dataset already

included, and its raw data are not openly redistributable. Including it would add a second instance of an assay type rather than a new assay technology. The Patricelli study[21] profiles only a small number of inhibitors against native kinases as a proof of concept for the KiNativ platform. A ranking-stability analysis requires many compounds to produce a meaningful ranking, so this dataset is not suited to the present comparison.

Selectivity Definitions and Parameterization

We implemented four selectivity definitions, each including parameters whose values are not standardized in the literature. We discuss each definition below. Throughout, a compound's binding profile is the vector of its affinities $\mathrm{p}K_{d,i}$ across the N panel kinases, with higher $\mathrm{p}K_d$ denoting tighter binding.

S-score. Karaman et al.[4] defined the selectivity score as the fraction of the panel that a compound binds: the number of kinases bound with a dissociation constant below a cutoff, divided by the number of distinct kinases tested. They reported $S(3\ \mu\mathrm{M})$ and $S(100\ \mathrm{nM})$, for K_d cutoffs of 3 μM and 100 nM. We implemented the same fraction, with the cutoff expressed as a threshold τ on $\mathrm{p}K_d$:

$$S(\tau) = \frac{1}{N} \sum_{i=1}^{N} \mathbf{1}[\mathrm{p}K_{d,i} > \tau]. \tag{1}$$

We used $\tau \in \{5.5, 5.75, \ldots, 8.0\}$ (11 values), spanning K_d cutoffs from approximately 3 μM to 10 nM, which includes both of Karaman et al.'s cutoffs ($\tau \approx 5.5$ and $\tau = 7$).

Selectivity entropy. Uitdehaag and Zaman[5] treat the compound's binding across the panel as a probability distribution and report its Shannon entropy. A profile concentrated on few kinases has low entropy and is selective. In the original definition, the weight on each kinase is its association constant $K_{a,i} = 1/K_{d,i}$ (the share of compound that would be bound to kinase i in an equimolar mixture), so $p_i = K_{a,i} / \sum_j K_{a,j}$. We implemented the entropy on the same log-affinity scale used by the other metrics, weighting each kinase by its affinity in

excess of a baseline β:

$$H(\beta) = -\sum_{i=1}^{N} p_i \log_2 p_i, \quad p_i = \frac{\max(\mathrm{p}K_{d,i} - \beta,\, 0)}{\sum_j \max(\mathrm{p}K_{d,j} - \beta,\, 0)}, \tag{2}$$

where β is the affinity below which binding is treated as negligible. This makes β an explicit parameter (its under-specification motivates property D3 below). We swept $\beta \in \{5.0, 5.25, \ldots, 6.5\}$ (7 values), a narrower range than the S-score's τ sweep above. Both parameters key off the same threshold-crossing event, so sweeping either one far enough leaves a growing share of compounds with no kinase above the cutoff (on Klaeger, 7.2% of compounds at $\mathrm{p}K_d = 6.0$ versus 55% at $\mathrm{p}K_d = 8.0$, `dataset_descriptive_stats.py`), degenerating the S-score toward an uninformative 0 and the entropy and Gini weights toward an undefined 0/0 alike. The τ range (whose S-score is not necessarily immune to this) is wider because it was chosen to include Karaman et al.'s published cutoffs (3 μM and 100 nM) and extend to 10 nM. β has no equivalent published reference points and was kept closer to the $\mathrm{p}K_d = 6$ active threshold used elsewhere in this work. The association-constant and log-affinity weightings are distinct reasonable choices that can rank compounds differently, which is itself an instance of the definitional instability this paper characterizes. We quantify the difference below.

Gini coefficient. Graczyk[6] adapted the Gini coefficient from economics: kinases are sorted by activity and the cumulative fraction of total binding is plotted against the cumulative fraction of kinases (a Lorenz curve). The Gini coefficient, which measures how unequally binding is distributed across the kinome, is twice the area between this curve and the diagonal of equal distribution. The original work computes it from percent-inhibition values. As with the entropy above, we implemented the Gini coefficient on the same log-affinity scale for the same reason. We applied the standard discrete Gini formula to that baseline-subtracted

log-affinity profile:

$$G(\beta) = \frac{2\sum_{i=1}^{N} i \cdot x_{(i)}}{N\sum_{i=1}^{N} x_{(i)}} - \frac{N+1}{N}, \qquad x_{(i)} = \max(\mathrm{p}K_{d,i} - \beta,\, 0) \text{ sorted ascending.} \tag{3}$$

We swept β over the same range as for the entropy. A compound whose binding is concentrated on one kinase has a highly unequal distribution (G is near 1), while a compound that binds many kinases equally has G near 0. Higher values therefore indicate greater selectivity.

Ratio. Classical pharmacological selectivity is pairwise. It quantifies how much more tightly a compound binds its primary target than an off-target. Karaman et al.[4] also reported a score of this kind alongside the S-score: the fraction of tested kinases bound within tenfold of the affinity for the compound's intended primary target, which combines a ratio cutoff with an S-score-style count. Because it anchors on the intended primary target rather than on the highest-affinity kinase, it does not depend on which of two near-tied top kinases is designated the primary target, the ambiguity behind the ratio's near-tie instability (Type 2, Results) and property D2 (below). It does, however, require an annotated primary target for every compound, whereas $R(k)$ needs only the profile itself. We implemented the pairwise comparison directly, as the fold-selectivity, in log units, between the primary target and the k-th ranked off-target:

$$R(k) = \mathrm{p}K_{d,(1)} - \mathrm{p}K_{d,(k+1)} = \log_{10} \frac{K_{d,(k+1)}}{K_{d,(1)}}, \tag{4}$$

where $\mathrm{p}K_{d,(j)}$ is the j-th largest affinity in the profile and $k \in \{1, 2, 3, 4, 5\}$ selects the reference off-target. We capped k at 5 because, on the sparser Klaeger panel, a large share of compounds already have fewer than five active kinases (median 5.5; 45% have fewer than five, `dataset_descriptive_stats.py`). For those compounds, a larger k would compare the primary target against the assay floor rather than a genuine off-target.

This floor fallback is a deliberate departure from Bosc et al.'s stated convention, under

which a ratio-based score is left undefined when a compound has no primary target or off-target to compare (Background); our implementation instead returns a numerical value for every compound by substituting the assay floor for a missing off-target rather than leaving the score undefined. We flag this explicitly because it means $R(k)$'s D1 violation (Table 3) is a property of this implementation choice, not of the definition as Bosc et al. state it: an implementation that followed their convention would instead satisfy D1 by construction, the same way the candidate measure does.

$R(k)$ takes the affinity gap to a single chosen competitor, the k-th ranked off-target. The field also assigns the concentration- and target-dependent selectivity (CATDS) score of Klaeger et al.[7] to this family, though we find that it does not behave like a member of it. CATDS divides each target's concentration-dependent binding reduction by the sum of binding reductions over *every* target in the panel, rather than comparing to a single chosen competitor, which is a distributional construction. We therefore represent the ratio family by $R(k)$ alone and treat CATDS separately in the Discussion (`bosc_catds_comparison.py`, Supporting Information Table S2). We use the simpler fold-selectivity because the reference off-target index k is then an explicit parameter whose effect can be studied directly. A larger affinity gap to the nearest competitor means greater selectivity, so higher $R(k)$ indicates greater selectivity.

Faithfulness to the original definitions. The S-score and the ratio are computed in the form given by their original sources (a panel-hit fraction and a primary-to-off-target affinity ratio, respectively). For the entropy and the Gini coefficient, we used the baseline-subtracted log-affinity profile in place of the original association constants (entropy) and percent-inhibition values (Gini), so that all four metrics operate on one common scale, and making baseline β an explicit calculation parameter. Every result reported for the entropy and the Gini coefficient elsewhere in this paper, including Table 1's two-family clustering, therefore pertains to this re-operationalized, baseline-subtracted form, not to the association-

constant or percent-inhibition forms of the original publications. The two operationalizations rank compounds differently enough to change which family entropy appears to belong to. To test whether this choice determines our conclusions, we recomputed the entropy and Gini with association-constant weighting and compared the resulting rankings with the log-affinity forms. For the entropy this is the original form; for the Gini, whose original input is percent inhibition, it is an affinity-scale analogue. In Figure 1 the association-constant forms are computed over every entry of the matrix, floor-imputed ones included; the robustness check below instead restricts them to measured entries. The two operationalizations differ substantially: for entropy, Spearman $r = 0.42$ (Davis) and -0.32 (Klaeger), while for Gini, $r = 0.11$ and -0.45 (Figure 1).

As a further robustness check, we recomputed the full cross-definition correlation matrix of Table 1 with entropy replaced by its association-constant form on Davis, Klaeger, and Metz. The Gini coefficient was replaced by its native percent-inhibition form (no baseline subtraction) on Anastassiadis, the one dataset where that native scale exists. On Klaeger and Metz, which have missing data, this association-constant form was computed per compound over only its measured targets, so that untested kinase–compound pairs are excluded from the weighting. The Davis dataset is a complete matrix with no untested pairs, so no such exclusion applies there.

On Davis, entropy decouples from the rest of the distribution family (S-score $r = 0.42$, Gini $r = 0.42$), close to its correlation with the ratio ($r = 0.34$), consistent with Figure 1. On Klaeger and Metz, entropy correlates more strongly with the ratio ($r = 0.79$ and 0.76) than with the S-score ($r = 0.57$ and 0.24) or the manuscript's own log-affinity Gini ($r = 0.69$ and 0.42). Under this operationalization entropy does not cluster with the distribution family at all on either sparse-coverage dataset. The native percent-inhibition Gini on Anastassiadis stays only partly consistent with the log-affinity form: it correlates with the entropy ($r = 0.63$), but weakly and about equally with the S-score ($r = 0.35$) and the ratio ($r = 0.31$), so the two-family pattern is not clearly preserved there either.

We draw two conclusions from this. First, the two-family structure we report throughout is a property of the operationalization we adopt, justified above for placing all four metrics on one common, parameter-explicit scale, and should not be read as a property of the named metrics "entropy" and "Gini" as originally published. Second, the direction of the effect is dataset-dependent (entropy decouples from the distribution family on Davis but is pulled toward the ratio family on Klaeger and Metz), so no single correction recovers one stable cross-dataset answer. The operationalization choice is itself an additional, uncharacterized source of definitional instability, on top of the three we characterize in Results. Full correlation matrices for all four datasets are reported in Supporting Information (Table S3) and in the code repository (`original_operationalization_robustness.py`).

Rank Stability Analysis

For each definition and each parameter value, compounds were ranked from most selective (rank 1) to least selective (rank n, where n is the number of compounds in the dataset). The S-score and the raw Shannon entropy both run opposite to this convention (a higher fraction of panel hits, or a higher raw entropy, indicates a *more* promiscuous compound), so both are negated before ranking throughout this paper. The Gini coefficient and $R(k)$ already increase with selectivity and are used unnegated. Ties were broken by average rank. The entropy and Gini are undefined for a compound with no kinase above the baseline β, because the normalizing sum is zero. Our implementation sets both scores to 0 in that case, which ranks the compound as maximally selective under the (negated) entropy, tied with single-target profiles, and as minimally selective under the Gini. A compound can be undefined at some baselines of the sweep and defined at others, so this convention enters its rank instability directly; Results reports the contrast under the alternative of computing instability only over configurations in which the score is defined. The ratio substitutes the assay floor for a missing off-target (above), and the S-score is defined (0) for every compound. This yielded 30 rank vectors per dataset for Davis, Klaeger, and Metz (11 for the S-score, 7 for entropy, 7

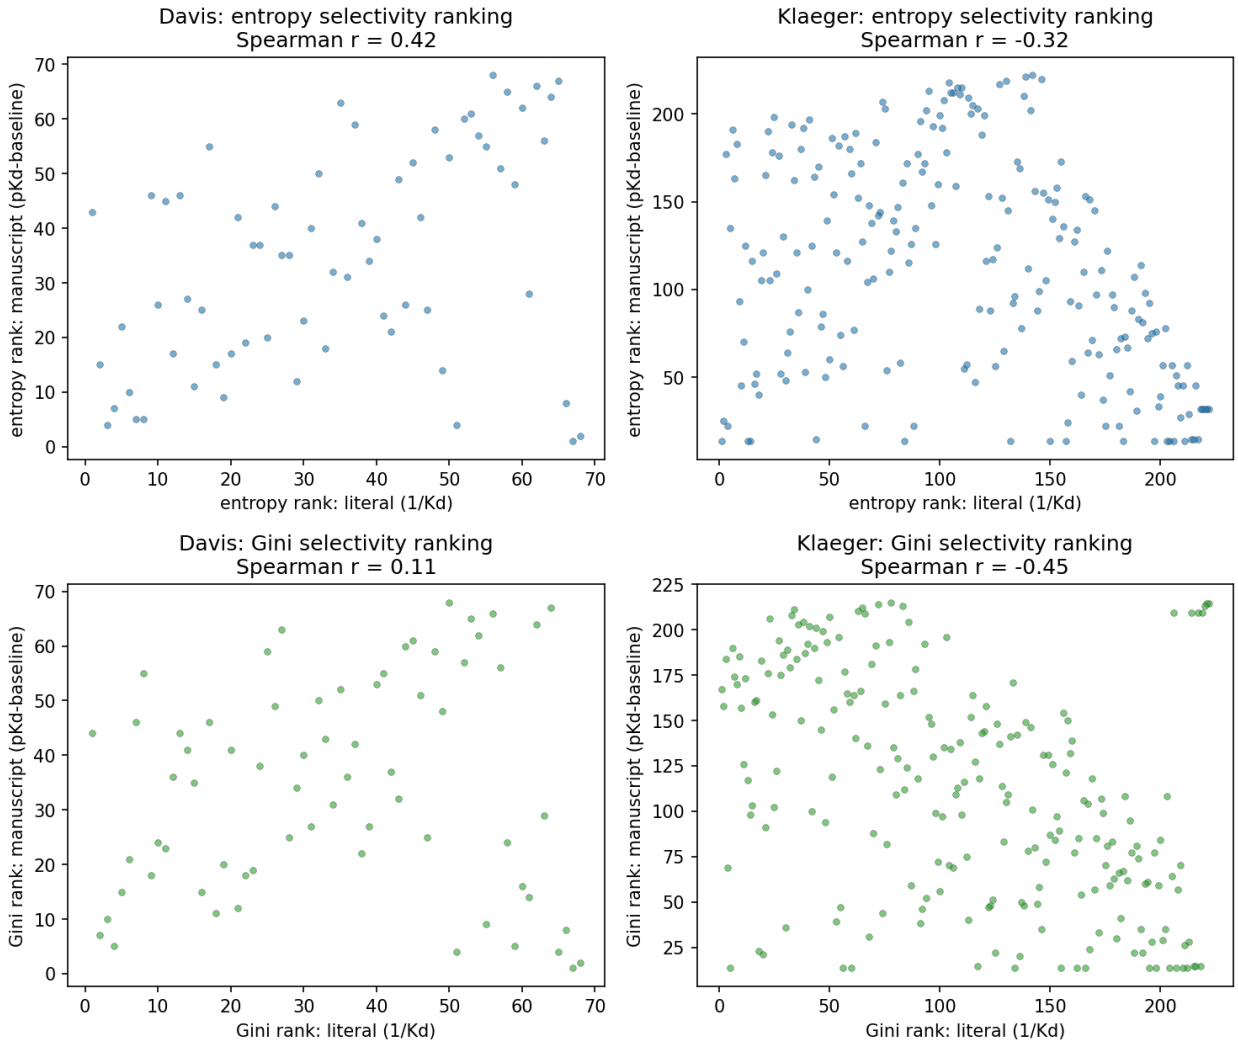


Figure 1: Operationalization sensitivity of the entropy and Gini coefficient. For each dataset, the compound selectivity ranking computed in a form close to the original publication (association-constant weighting) is plotted against the baseline-subtracted log-affinity form used in the main text. Top row: selectivity entropy; bottom row: Gini coefficient. The two operationalizations of the same named metric differ substantially, and for entropy the relationship between them even changes sign on Klaeger ($r = 0.42$ Davis, -0.32 Klaeger), confirming that the operationalization of a named metric is a source of the definitional instability.

for Gini, and 5 for ratio), and 28 for Anastassiadis, whose coarser native percent-inhibition scale supports only 9 S-score threshold values rather than 11.

For each compound, rank instability was quantified as the standard deviation of its rank across all configurations for that dataset (overall instability) and separately across parameter configurations within each definition family (within-family instability). Cross-definition-family agreement was assessed using Spearman rank correlation between the median rank vectors of each definition family. The median was taken across parameter values within each family, rather than the mean, because the median is less sensitive to the inclusion or exclusion of outliers than the mean.

Associations between rank instability and binding profile features were assessed using Spearman correlation, restricted to the Klaeger and Davis datasets. These are the two datasets on a pK_d scale, so the same activity threshold and affinity-based features apply directly, allowing the predictor structure to be checked for replication across a second dataset. Profile features computed for each compound included: the number of active kinases (p$K_d > 6$), the gap between the top and second-ranked affinities, the standard deviation and range of active affinities, the absolute affinity of the second-ranked target, and the skewness of the active affinities. All correlations are reported with two-tailed p-values. Significance thresholds of $p < 0.05$, $p < 0.01$, and $p < 0.001$ are denoted *, **, and *** respectively. Analyses were performed in Python 3.10 using NumPy, pandas, and SciPy.[22–24] A secondary analysis relating selectivity rankings to clinical adverse-event outcomes (FAERS report rates and FDA-label discontinuation rates for the FDA-approved Klaeger compounds) is summarized in the Limitations section. The analysis code and full per-definition results are available in the code repository (see Data and Software Availability).

Motivated Properties

The instability patterns reported in Results motivate four properties. Each is stated formally below, together with the failure mode it targets and how a measure is assessed against it;

Results evaluates the four definitions and a candidate measure against them (Table 3). Let $\mathbf{x} = (x_1, \ldots, x_N) \in \mathbb{R}^N$ denote a compound's binding profile across N kinases, where $x_i = \mathrm{p}K_{d,i}$ and N is the panel size, which varies across compounds within a dataset (Anastassiadis, Metz) and across datasets. A selectivity measure is a function $f : \mathbb{R}^N \to \mathbb{R}$, defined for every panel size N, where higher values indicate greater selectivity. A ranking over a set of compounds is defined by ordering their f-scores. We write x_0 for the assay's detection floor and τ^* for the activity threshold ($\mathrm{p}K_d = 6$ on the $\mathrm{p}K_d$ datasets).

D1 (reliability threshold). There exists $\tau^* > 0$ such that $f(\mathbf{x})$ is undefined, or explicitly flagged as unreliable, when $\max_i x_i \leq \tau^*$. A compound with no binding above the activity threshold has no target on which to ground a selectivity judgment, and its rank is otherwise set by whatever convention a measure uses for such profiles (Type 1, Results). D1 is a reporting rule, not a property of a functional form, and is assessed as such.

D2 (bounded gap sensitivity). For any $\mathbf{x}$ satisfying D1, if f depends on a choice of reference off-target indexed by k, there must exist g with $g(0) = 0$ such that $|f(\mathbf{x}; k) - f(\mathbf{x}; k+1)| \leq g(x_{(k)} - x_{(k+1)})$. When the k-th and $(k+1)$-th ranked affinities are nearly equal, it is ambiguous whether the $(k+1)$-th kinase is a target or the first off-target; moving the reference off-target past it (from k to $k+1$) must then change the score by only a small amount. This targets the ratio's near-tie instability (Type 2, Results) and is assessed from each measure's functional form. A measure with no reference-off-target parameter satisfies D2 vacuously.

D3 (distributional consistency). Let $B = [x_0, \tau^*]$, the interval between the assay's detection floor and the activity threshold, which contains every defensible placement of an activity boundary. Let θ be a parameter of f that an analyst must choose on the affinity scale (the S-score's τ, the entropy's and Gini's β). A pair of compounds $\mathbf{x}, \mathbf{y}$ *reverses in* B if $f(\mathbf{x}; \theta) - f(\mathbf{y}; \theta)$ is positive at some $\theta \in B$ and negative at another. D3 requires that at most a fraction ε of the pairs of compounds that D1 allows to be scored reverse in B. We use $\varepsilon = 0.01$; every verdict in Table 3 is the same for any ε from 1% to 7% (Results). This

targets the baseline sensitivity of compounds with few active kinases (Results) and is assessed empirically, by computing the rankings across B on each dataset. A measure with no such parameter satisfies D3 vacuously.

D4 (monotonicity under weak off-target addition). If $\mathbf{x}' = (x_1, \ldots, x_N, x_{N+1})$ with $x_{N+1} \leq \tau^*$, then $f(\mathbf{x}') \leq f(\mathbf{x})$. That is, adding a sub-threshold binding interaction must not increase a compound's score; otherwise a compound can be made to look more selective by profiling it against more kinases it barely binds, which matters whenever panels differ in size or composition. D4 is assessed from each measure's functional form.

A Candidate Measure

The candidate measure is introduced to show that the four properties can be met together, not to outperform existing measures on these data. For a profile $\mathbf{x}$ with $\max_i x_i > \tau^*$ (undefined otherwise, satisfying D1), define

$$w_i = \mathbf{1}[x_i > x_0]\, T \log\!\big(1 + e^{(x_i - \beta)/T}\big), \qquad p_i = \frac{w_i}{\sum_j w_j}, \qquad f(\mathbf{x}) = \sum_i p_i \log_2 p_i, \tag{5}$$

where x_0 is a constant representing the assay detection floor ($\mathrm{p}K_d = 5.0$ on Davis and Klaeger, $\mathrm{p}K_i = 4.0$ on Metz, and a 20% inhibition noise level on Anastassiadis), β is a smooth emphasis baseline (set to x_0 in the results below), and T is a smoothing width in the native units of the profile. We set $T = \tau^* - x_0$, the distance from the floor to the reliability threshold: 1 $\mathrm{p}K_d$ unit on Davis and Klaeger, 2 $\mathrm{p}K_i$ units on Metz, and 30 percentage points on Anastassiadis, whose native scale is 0–100% inhibition rather than log-affinity. With $\beta = x_0$ the measure then depends on a profile only through $(x_i - x_0)/(\tau^* - x_0)$, so it does not depend on the units of the readout. The function f is the negative Shannon entropy of the smoothed, normalized profile, i.e. the negative logarithm of the effective number of targets the compound binds.

The candidate is a gated, smooth-hinge variant of the selectivity entropy and differs from it in three ways, one for each property the entropy fails: (1) an explicit reliability gate, for

D1; (2) a non-binder gate fixed at the floor x_0, so that no free parameter sets the activity boundary, for D3; and (3) a smooth softplus hinge above x_0 in place of the entropy's hard cutoff $\max(x_i - \beta, 0)$, for D4. The hard cutoff gives a kinase just above β an arbitrarily small weight, which is what lets a sub-threshold addition raise the entropy's score; the softplus gives every gated kinase a weight of at least $T \ln 2$, which rules this out:

Proposition. With $\beta = x_0$ and $T \geq (\tau^* - x_0)/\ln(2^e - 1) \approx 0.58\,(\tau^* - x_0)$, adding a kinase with affinity at most τ^* never raises the candidate's score.

Proof. A kinase at or below x_0 has zero weight and leaves the score unchanged. Otherwise every gated kinase, the added one included, has weight above $m = T \ln 2$, and the added one has weight $a \leq T \ln(1 + e^{(\tau^* - x_0)/T}) \leq e\,m$, the last step by the condition on T. Let the original profile have n gated kinases with total weight $W > nm$ and entropy H (in nats), and let $q = a/(W + a)$. After the addition the entropy is $(1 - q)H + h(q)$, where h is the binary entropy, so it does not decrease if and only if $H \leq h(q)/q$. Because h is concave with $h(0) = 0$, $h(q)/q$ decreases in q. Since $q < r/(n + r)$ with $r = a/m \leq e$, this gives $h(q)/q > \ln \frac{n+r}{r} + \frac{n}{r} \ln \frac{n+r}{n}$. As $H \leq \ln n$, it suffices that $\ln n$ is at most this bound, which rearranges to $\ln r \leq (1 + \frac{n}{r}) \ln(1 + \frac{r}{n})$. The right-hand side is at least 1 for all $n, r > 0$, because $\ln(1 + t) \geq t/(1 + t)$, and $\ln r \leq 1$. □

The choice $T = \tau^* - x_0$ satisfies this condition on every dataset.

Results

Selectivity Definitions Cluster into Two Families

Across 30 parameter configurations (11 S-score thresholds, 7 entropy baselines, 7 Gini baselines, 5 ratio off-target indices), compound rankings were computed for all drugs in both datasets. Spearman correlations between the median ranking of each definition family reveal a consistent two-cluster structure (Table 1). On the Davis dataset, S-score, entropy, and Gini are mutually correlated ($r = 0.946$–0.999), while ratio correlates substantially less with

each ($r = 0.34$–0.46). On the larger Klaeger dataset, the same structure is preserved but with lower within-cluster correlations ($r = 0.72$–0.90), reflecting greater diversity of binding profiles among the 222 clinical compounds. The ratio definition is the consistent outlier in both datasets, with cross-family correlations of $r = 0.22$–0.62.

Table 1: Spearman correlations between median rankings by definition family. Upper triangle: Davis ($n = 68$). Lower triangle: Klaeger ($n = 222$).

	S-score	Entropy	Gini	Ratio
S-score	—	0.946	0.952	0.456
Entropy	0.901	—	0.999	0.343
Gini	0.722	0.866	—	0.343
Ratio	0.223	0.483	0.616	—

This two-family structure reflects a fundamental difference in what the definitions measure. The S-score, entropy, and Gini all quantify the concentration of activity across the full kinase panel. They differ in functional form but respond to the same underlying property of the binding profile. The ratio definition, by contrast, quantifies the gap between the primary target and a specific off-target, ignoring the remainder of the binding profile entirely. A compound that strongly inhibits one kinase and moderately inhibits twenty others will be ranked as highly selective by ratio (large gap to the second target) but as moderately promiscuous by entropy and Gini (activity distributed across many targets).

The Two-Family Structure Replicates Across Assay Technologies

The Davis and Klaeger pattern of Table 1 extends to the two further datasets, across three assay technologies and both readout types (Figure 2). In all four, the S-score, entropy, and Gini remain mutually correlated (within-family Spearman $r = 0.72$–0.999) while the ratio definition is the consistent outlier. Its correlation with the distribution-based family is 0.49–0.56 on the Anastassiadis functional-activity dataset and only 0.14–0.19 on the Metz pK_i dataset. The degree of separation varies with the assay, but in every case the ratio definition measures something the distribution-based definitions do not.

The instability findings also replicate. Compounds with no active kinases (Type 1, below) are markedly more rank-unstable than active compounds wherever such compounds exist: mean rank $\sigma = 36.8$ vs. 21.7 on Anastassiadis and 178.5 vs. 108.1 on Metz, matching the 72.9 vs. 31.3 contrast on Klaeger (Davis contains no zero-active compounds). As on Klaeger, this contrast reflects the convention used for undefined entropy and Gini scores (Methods; Type 1, below). Within the distribution-based family, the number of active kinases negatively predicts entropy instability in every dataset ($r = -0.28$ to -0.43). The top_1–top_2 affinity gap predicts ratio instability on three of the four datasets, Davis, Klaeger and Metz ($r = -0.44$, -0.34, -0.27; all $p < 0.001$), but not on Anastassiadis ($r = -0.104$, $p = 0.23$), as detailed below.

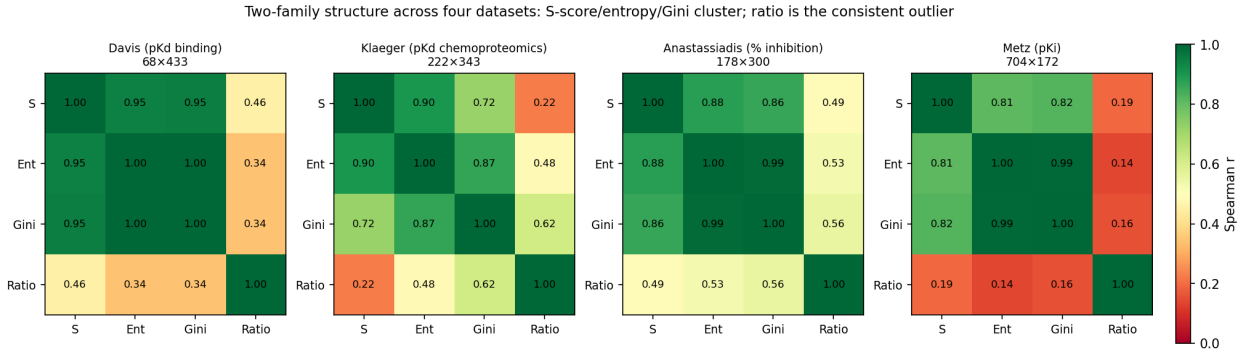


Figure 2: Spearman correlations between the median rankings of the four selectivity definitions, computed independently on each of the four datasets (compound × kinase panel sizes shown). In every dataset the S-score (S), entropy (Ent), and Gini cluster together while the ratio is the consistent outlier (the ratio row and column are the least green throughout), confirming the two-family structure across competition-binding, chemoproteomic, and functional-activity assays.

Rank Instability Is Structured and Associated with Profile Shape

Overall rank instability (measured as the standard deviation of a compound's rank across all 30 configurations) varies substantially across compounds (Klaeger: mean $\sigma = 34.3$, range 7.0–90.7; Davis: mean $\sigma = 8.5$, range 1.8–19.3). Instability is not uniformly distributed: within each definition family, specific binding profile features are associated with which

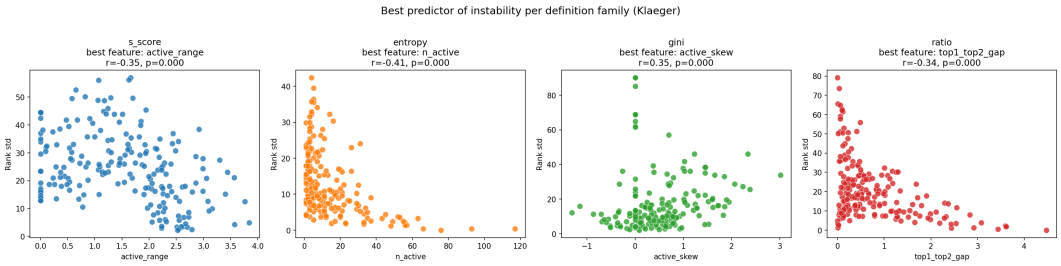


Figure 3: Best predictor of rank instability within each definition family (Klaeger dataset, active compounds, $n = 206$). Ratio instability is best predicted by the top_1–top_2 affinity gap; entropy, Gini, and S-score instability are best predicted by features of the active affinity distribution.

compounds receive unstable rankings (Table 2).

Within the ratio family, the gap between the top and second-ranked affinities (top_1–top_2 gap) is the strongest predictor of instability (Figure 3) ($r = -0.444$, $p < 0.001$ on Davis, and $r = -0.341$, $p < 0.001$ on Klaeger). Compounds with near-identical affinities for their two best-bound kinases receive highly variable ratio rankings because the choice of reference off-target ($k = 1$ vs. $k = 2, \ldots, 5$) determines whether the compound appears selective or promiscuous.

Within the entropy family, the number of active kinases ($\text{p}K_d > 6$) negatively predicts instability ($r = -0.279$, $p < 0.05$ on Davis, $r = -0.409$, $p < 0.001$ on Klaeger; for the Gini, see below). Compounds with many active kinases have well-defined broad profiles whose entropy and Gini scores are stable across baseline choices, while compounds with few active kinases are sensitive to the baseline because small changes can move targets in or out of the active set, altering the distribution substantially. The active affinity range is also a significant predictor for the S-score and entropy on Klaeger ($r = -0.35$ and -0.27, both $p < 0.001$). Gini is not significantly predicted by this feature ($r = -0.10$, $p = 0.14$). It is consistent with the interpretation that heterogeneous profiles (whose affinities span a wide range) yield more reproducible rankings across parameter choices than flat profiles (in which many near-equal affinities reshuffle as the parameter changes).

Gini instability's strongest predictor on Klaeger ($r = +0.346$, $p < 0.001$) is skewness,

comparable in strength to the S-score's dependence on active affinity range ($r = -0.354$). Unlike the other features above, it targets what Gini measures directly, i.e. the shape of the active-affinity distribution's inequality. The positive sign of this correlation is consistent with right-skewed active distributions being the least stable. Most active kinases bound only modestly above threshold, with one (or a few) much more potent outliers, concentrating mass near the baseline boundary. A small shift in β moves that mass in or out of the active set. Left-skewed profiles (most active kinases tightly bound, few sitting near threshold) are comparatively insensitive to that choice. As with active affinity std/range above, the relationship does not replicate on Davis ($r = +0.043$, $p = 0.73$).

Unlike the other families, the S-score's instability is significantly predicted by every profile-shape feature except the top_1–top_2 gap (Table 2), and most strongly by the spread of active affinities. Compounds whose active affinities cluster near the threshold gain or lose panel hits under small shifts in τ, whereas those with a wide affinity range retain a stable hit count regardless of the exact cutoff. This S-score pattern is specific to Klaeger: on Davis, however, no profile feature significantly predicts S-score instability (all $|r| \leq 0.16$; Table 2, right). A possible explanation is that a threshold count's instability is governed by where each assay's detection window falls relative to the cutoff, and the two assays place that window very differently: 3.6% of the Klaeger matrix is active versus 17.9% of Davis (Methods). Two relations in Table 2 are significant on both pK_d datasets: the ratio family's dependence on the top_1–top_2 gap ($p < 0.001$ on both), and the entropy's dependence on the number of active kinases. The ratio-gap pairing is the more consistent of the two in magnitude ($r = -0.444$ Davis, -0.341 Klaeger), while the entropy pairing holds on both but is considerably weaker on the smaller Davis panel ($r = -0.409$, $p < 0.001$, on Klaeger vs. -0.279, $p < 0.05$, on Davis). The ratio-gap dependence, which holds on Metz but fails on Anastassiadis, replicates on three of the four datasets, and specifically not on the percent-inhibition readout. Gini shows the reverse pattern in strength to the entropy, reaching significance on Davis ($r = -0.300$), but not on Klaeger ($r = -0.093$).

Table 2: Spearman correlations between binding profile features and within-family rank instability, on the Klaeger ($n = 206$ active drugs) and Davis ($n = 68$) datasets. Columns are the S-score (S), entropy (E), Gini (G), and ratio (R) families. The ratio family's dependence on the top_1–top_2 gap replicates strongly on both datasets. The S-score's dependence on the active-affinity spread does not (no feature is significant on Davis), and neither does Gini's dependence on active-affinity skew. Significance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

	Klaeger ($n = 206$)				Davis ($n = 68$)			
Feature	S	E	G	R	S	E	G	R
n_{active}	-0.272^{***}	-0.409^{***}	-0.093	-0.040	-0.153	-0.279^{*}	-0.300^{*}	$+0.294^{*}$
top_1–top_2 gap	$+0.039$	$+0.082$	$+0.031$	-0.341^{***}	-0.079	-0.007	-0.029	-0.444^{***}
Active affinity std	-0.281^{***}	-0.170^{*}	-0.117	-0.107	$+0.009$	-0.015	$+0.005$	$+0.040$
Active affinity range	-0.354^{***}	-0.273^{***}	-0.102	-0.116	$+0.011$	-0.111	-0.098	$+0.134$
top_2 pK_d	-0.314^{***}	-0.234^{***}	-0.192^{**}	$+0.186^{**}$	-0.055	-0.203	-0.187	$+0.296^{*}$
Active affinity skew	-0.162^{*}	-0.045	$+0.346^{***}$	-0.108	$+0.026$	$+0.066$	$+0.043$	$+0.069$

Skewness is undefined for compounds with fewer than three active kinases, and the standard deviation and range for compounds with fewer than two; Table 2 assigns these compounds the value 0 (the vertical clusters at zero in Figure 3). Recomputing each correlation over only the compounds for which the feature is defined (161 and 180 of the 206 active Klaeger compounds, respectively) strengthens every Klaeger association reported above, including Gini–skewness ($r = +0.482$) and the S-score's and entropy's dependence on active affinity range ($r = -0.479$ and -0.401; all $p < 0.001$). It also makes the Gini's and the ratio's weak dependence on active-affinity spread nominally significant ($|r| = 0.18$–0.28, $p < 0.05$). On Davis, where only three compounds have fewer than three active kinases, no correlation changes materially (`feature_placeholder_sensitivity.py`). The reported associations therefore do not depend on the placeholder.

Five additional candidate features were tested: two concentration measures from industrial-organization economics ((1) the top-target concentration ratio (CR_1) and (2) the Herfindahl–Hirschman Index (HHI)). These were sourced from the same field as the Gini coefficient.[6] We also tested (3) the raw primary-target affinity (top_1 pK_d), (4) a near-tie count generalizing the top_1–top_2 gap, and (5) the active-affinity coefficient of variation. CR_1 and HHI turned out to be near-collinear with n_{active}, which is already in the table ($r = -0.98$ to -0.99 and -0.998, respectively, on both datasets), and thus they carry essentially the same information.

The remaining three were weaker predictors than the features already included, in every family and both datasets (full numbers in the code repository).

Three Mechanistically Distinct Sources of Instability

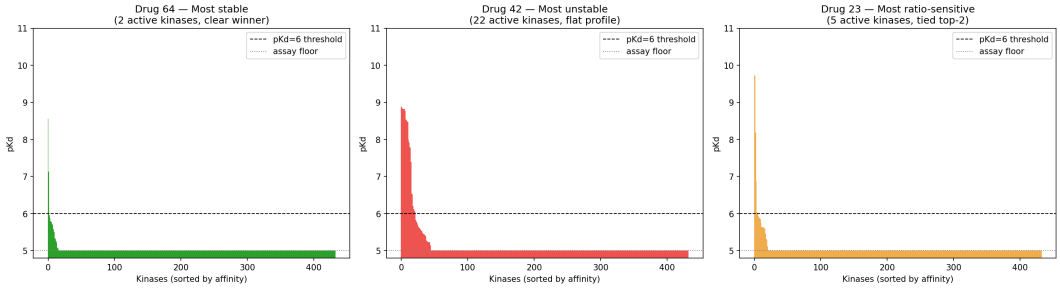


Figure 4: Binding profiles of the most stable, most unstable, and most ratio-sensitive Davis compounds. Davis has no zero-active (Type 1) compounds; the right panel illustrates Type 2. Left: Drug 64, most stable ($\sigma = 1.8$), two active kinases. Center: Drug 42, most unstable ($\sigma = 19.3$), 22 active kinases with no dominant target. Right: Drug 23, most ratio-sensitive, five active kinases including two near-tied top targets ($\Delta \mathrm{p}K_d = 0.023$). Dashed line: $\mathrm{p}K_d = 6$ activity threshold; dotted line: $\mathrm{p}K_d = 5$ assay floor.

Examination of the most and least unstable compounds in the Klaeger dataset reveals three structurally distinct sources of rank instability, each implicating a different aspect of how the definitions are constructed. Figure 4 shows example profiles from Davis, where the extremes are easier to see on the complete matrix.

Type 1: Zero-active compounds. The five most unstable Klaeger compounds (AZD-8055, BMS-911543, Vatalanib, Filgotinib, ARRY-380; rank $\sigma > 76$) all have zero active kinases at the $\mathrm{p}K_d > 6$ threshold in the Klaeger panel. This is a property of the assay result, not of the compounds' known pharmacology: all five are established, potent inhibitors of specific kinase targets (mTOR, JAK2, VEGFR2, JAK1, and HER2, respectively), none of which were captured as high-confidence sub-micromolar hits in the Klaeger chemoproteomic screen. A plausible explanation (which would require additional data to confirm) is low expression or poor capture of those targets in the profiled cell lysates.

Within the Klaeger dataset, all four definitions assign these compounds numerical selec-

tivity scores based only on sub-threshold values, with no primary target on which to ground a selectivity judgment. For entropy, Gini, and the ratio, the resulting rankings are arbitrary: different parameterizations produce different orderings of compounds that are, in practice, indistinguishable from one another. The S-score is the exception because it evaluates to exactly 0 for every zero-active compound regardless of threshold, so these compounds are tied rather than reordered under the S-score, giving a stable but degenerate ranking. None of the four definitions withholds a score from these compounds, so all four fail D1 (Table 3, below).

Mean rank instability for zero-active compounds is $\sigma \approx 73$, compared with $\sigma \approx 31$ for compounds with at least one active kinase ($p < 0.001$, Mann–Whitney U test). This contrast depends on how the undefined entropy and Gini scores of these compounds are ranked (Methods): computed only over the configurations in which each compound's score is defined, zero-active compounds are no less stable than active ones (D1, below). What sets them apart is that their ranks are fixed by a convention rather than by measured binding. Davis has no such compounds at all: every Davis compound has at least one active kinase, consistent with Davis's fixed-concentration screening design, which reports a K_d only when activity is first detected at 10 μM. The same zero-active instability contrast reported here for Klaeger also holds on Anastassiadis and Metz (see the replication across assay technologies above).

Type 2: Near-tied top targets. When an active compound's two best-bound kinases have near-identical affinities, the ratio is highly sensitive to which is designated the primary target and which the reference off-target, because a small parameter change can swap their roles (the top_1–top_2 gap is the strongest predictor of ratio instability, above). Drug 23 in the Davis dataset (top_1–top_2 gap = 0.023 pKd units) is the extreme case: its two best-bound kinases are essentially indistinguishable, so it receives maximally unstable ratio rankings.

Type 3: Broad profiles with a dominant primary target. Among active compounds, the largest entropy-ratio disagreements in the Klaeger dataset arise for compounds with high n_{active} and a large top_1–top_2 gap (e.g. Dovitinib: $n_{\text{active}} = 26$, gap = 1.935; BMS-690514: $n_{\text{active}} = 20$, gap = 1.985). These compounds have one clearly dominant target but also bind

many kinases at moderate affinity. The ratio definition scores them as highly selective (large gap between primary and secondary target), while the entropy and Gini definitions score them as moderately promiscuous (activity spread across many targets). Neither characterization is incorrect, as they reflect genuinely different questions about selectivity.

Conversely, the most stable Klaeger compounds (XL-228, $\sigma = 7.0$; PF-03814735, $\sigma = 7.8$) are highly promiscuous, with 55–57 active kinases. This is because when a compound binds a large share of the profiled kinome, all definitions rank it as non-selective regardless of parameterization, producing low instability. Stability can also arise from the opposite extreme. The most stable Davis compound (Drug 64, Figure 4, left; $\sigma = 1.8$) has only two active kinases, but its affinity for them sits far enough above every other kinase's that no parameter choice in our sweep moves it off the top of the ranking. Both routes to stability, near-complete promiscuity and a sharply decisive, narrow profile, leave little room for parameter choice to change the outcome. The intermediate, ambiguous profiles described above are, on the other hand, consistently reported as unstable.

Panel Size Dependence of Selectivity Rankings

Karaman et al.[4] scored each compound's S-score on 500 random subpanels of each size and found that it varies more as panels shrink, concluding that a single small panel is unlikely to rank compounds reliably by overall selectivity. We extend this to the rankings themselves, for all four definitions, with a subsampling analysis on the two datasets with genuine sparse kinase coverage: Klaeger (343 kinases, chemoproteomic dose–response profiling) and Metz (172 kinases, a medicinal-chemistry campaign), in each of which compounds were tested against only a subset of the panel. The Davis and Anastassiadis datasets were not used for this analysis, because both are dense complete-matrix screens (every compound tested against every kinase by construction) rather than the sparse coverage of real profiling panels. Subsampling their kinases would also change the apparent number of active compounds in ways that confound the panel size effect with the zero-active instability characterized in the

instability taxonomy above.

For each panel size (Klaeger $p \in \{50, 80, \ldots, 320, 343\}$; Metz $p \in \{50, 80, \ldots, 170, 172\}$), we drew 50 random subsets of p kinases from the full panel, computed each selectivity definition on the subpanel, and measured agreement with the full-panel ranking using Spearman correlation. We call a panel size "recovering" the full-panel ranking once this mean correlation exceeds 0.90, adopting the same numerical bar that is conventionally sufficient to call a measurement "excellent".[25] Altering the choice of correlation cutoff changes the reported p^* values, however, the ratio definition remains the slowest to recover for any bar between 0.80 and 0.95, and the Gini coefficient the second-slowest on Klaeger (on Metz it ties the entropy at every bar). Only the relative order of entropy and the S-score is bar-dependent within this range (full numbers in Supporting Information, Table S1, and in the code repository).

On Klaeger, the four definitions differ substantially in their panel size requirements (Figure 5). Selectivity entropy and the S-score both reach mean $r = 0.90$ against the full-panel ranking at approximately 110 kinases. The Gini coefficient requires approximately 290 kinases to reach the same threshold. The ratio definition does not reach $r = 0.90$ at any subpanel size short of the full panel: it achieves only $r = 0.888$ even at 320 kinases, just 23 below the full panel of 343, and attains $r = 0.90$ only on the complete panel.

This result has a direct practical implication. Commercial kinase selectivity panels vary widely in size, from focused subpanels of a few dozen kinases to near-kinome panels of several hundred. On a 50-kinase subpanel, entropy- and S-score-based rankings are reasonably stable ($r \approx 0.78$–0.82), while Gini and ratio rankings are effectively uncorrelated with what would be obtained on a full kinome panel ($r = -0.09$ for Gini and $r = -0.08$ for ratio at $p = 50$). The practical consequence is that the smaller the panel, the more the choice of definition determines the answer. Ratio rankings in particular are unreliable at any panel size commonly used in practice, because the identity of the reference off-target (k-th ranked kinase) changes substantially as panel composition changes.

To test whether these patterns generalize beyond a single assay, we repeated the analysis on Metz (172 kinases; Figure 5, right). The distribution-based definitions recover the full-panel ranking from a small fraction of the panel (S-score at approximately 50 kinases, entropy and Gini at approximately 80), while the ratio again lags, reaching $r = 0.90$ only at 170 of the 172 kinases. The ratio's poor recoverability therefore holds on both sparse datasets, consistent with it measuring a distinct, panel-composition-dependent property, whereas the Gini coefficient's slow Klaeger recovery above does not recur on Metz, where it matches the entropy. Entropy and the S-score are fast on both datasets.

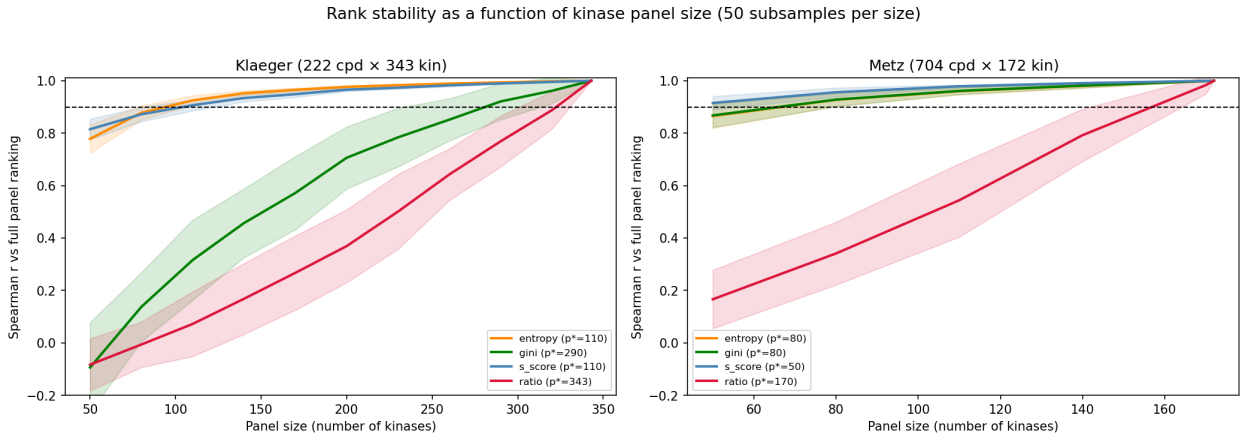


Figure 5: Mean Spearman correlation between rankings computed on random kinase subpanels and rankings on the full panel, as a function of subpanel size, for the two sparse-coverage datasets: Klaeger (left, 343 kinases) and Metz (right, 172 kinases). Shaded bands indicate ± 1 standard deviation across 50 subsamples; the dashed line marks $r = 0.90$.

Existing Definitions Against the Properties

Table 3 summarizes how the four definitions and the candidate measure fare against the properties D1–D4 defined in Methods. None of the four definitions satisfies all four. D1 is the only property all four violate, and the entropy comes closest overall: it satisfies D2 (vacuously) and fails D1, D3, and D4.

D1. None of the four definitions has a reliability condition, so each assigns a score to the 16 Klaeger compounds with no binding above the activity threshold. Their ranks are set by

Table 3: Satisfaction of the four motivated properties by existing selectivity definitions and by the candidate measure. ✓: satisfied substantively. (✓): satisfied vacuously, because the measure has no parameter of the kind the property constrains (no reference off-target for D2; no baseline for the ratio's D3). ✓†: satisfied only with the activity boundary fixed at the assay floor x_0; sweeping x_0 itself over B moves the ranking as much as the entropy's baseline does. ×: violated. Entries in the D3 row are determined empirically, on all four datasets; the other rows are assessed from each measure's functional form.

Property	S-score	Entropy	Gini	Ratio	Candidate
D1 (reliability threshold)	×	×	×	×	✓
D2 (bounded gap sensitivity)	(✓)	(✓)	(✓)	×	(✓)
D3 (distributional consistency)	×	×	×	(✓)	✓†
D4 (monotonicity, weak off-target)	×	×	×	✓	✓

convention rather than by measured binding. Pooled across the four definitions' parameter sweeps, their rank instability is $\sigma = 72.9$ against 31.3 for active compounds ($p < 0.001$, Mann–Whitney U test; Figure 6, left). The contrast is driven by Gini ($\sigma = 86.1$ versus 15.9) and entropy (17.1 versus 12.5), is marginal for the ratio (24.0 versus 20.7), and reverses for the S-score (18.5 versus 24.3), whose score is exactly 0 for every such compound. It is largely a product of the convention for undefined entropy and Gini scores (Methods): a zero-active compound's entropy and Gini are defined at only 1–4 of the 7 baselines swept (mean 3.2), and computed only over each compound's defined configurations, zero-active compounds are slightly *more* stable than active ones (entropy: $\sigma = 9.8$ vs. 12.1, $p = 0.085$; Gini: $\sigma = 9.5$ vs. 13.0, $p = 0.028$; 14 of the 16 compounds, the other two being defined at fewer than two baselines; `undefined_score_convention.py`). The case for D1 is therefore not that these scores are noisier, but that they rest on an arbitrary choice that D1 removes. Any of the four definitions could satisfy D1 by withholding or flagging such scores; for the ratio, our implementation's floor fallback is what violates it (Methods, "Ratio").

D2. The ratio violates D2 when its top two affinities are nearly tied. On the 17 Klaeger compounds with a top_1–top_2 gap below 0.05, mean $R(1) = 0.02$ (the near-tied second kinase is treated as the off-target) against mean $R(2) = 0.42$ (it is treated as a co-target; `d2_near_tie_sensitivity.py`). This is the Type 2 instability above, where the top_1–top_2

gap is the strongest predictor of ratio instability ($r = -0.341$, $p < 0.001$). The S-score, entropy, Gini, and candidate have no reference off-target and satisfy D2 vacuously; the S-score's sensitivity to kinases near its threshold is a D3 matter.

D3. Swept over $B = [x_0, \tau^*]$, the S-score, entropy, and Gini each reverse the order of 8–23% of pairs of scored compounds, on every dataset (S-score 8.0–20.6%, entropy 8.8–22.6%, Gini 9.1–22.4%; highest on Metz; `d3_pair_reversals.py`), far above $\varepsilon = 1\%$. The candidate's emphasis baseline reverses at most 0.72%, so every verdict in Table 3 is the same for any ε from 1% to 7%. This is consistent with the n_{active}–instability correlation of Table 2: compounds with few active kinases change their active sets as the boundary moves. Making β an explicit parameter is what makes this testable: the published originals leave it implicit ($K_a \to 0$ for the entropy, 0% inhibition for the Gini). The ratio has no baseline and satisfies D3 vacuously.

D4. The S-score and Gini violate D4 by construction: both are computed over the full panel of size N, so adding a kinase the compound does not bind lowers the S-score's hit fraction and raises the Gini's inequality, inflating the apparent selectivity. A compound can therefore be made to look more selective simply by profiling it against more inactive kinases, which matters whenever panels differ in size or composition, as between Davis and Klaeger. The ratio satisfies D4, because a sub-threshold kinase falls at or below its reference off-target. The entropy satisfies D4 only for additions at or below its baseline β, which carry zero weight; a kinase between β and τ^* carries positive weight and can raise its score ("Why the Smooth Hinge Matters", below). These arguments refer to the selectivity-oriented f (Methods), not to the raw entropy or S-score, which run the opposite way.

The Candidate Measure

The candidate measure (Methods) is evaluated here to show that the four properties can be met together. It is correctly oriented: its score falls with the number of active kinases on every dataset (Spearman -0.785 to -0.969). The non-binder gate is what makes this so. Without it, the non-binding kinases at the floor dominate the profile and the score inverts,

rising with the number of active kinases ($r = +0.94$ on Klaeger).

Because the candidate is a local modification of the entropy, the two rank the compounds they both score almost identically (full-panel Spearman +0.999 on Klaeger and +0.995 on Metz). The properties do not act by reordering compounds in the body of the distribution. They act on which compounds receive a score at all (D1, which flags 16 of 222 Klaeger compounds), on whether the ranking survives a shift in an analyst's parameter (D3), and on profiles that these datasets rarely or never contain (D4). The case for the properties therefore does not rest on these data. They are requirements on every binding profile, not only on the profiles four datasets happen to contain, and we propose them as checks that any new selectivity measure should pass. "Why the Smooth Hinge Matters", below, gives a concrete case in which the entropy fails one of them on a real dataset and the candidate provably does not.

The candidate satisfies D1 by its reliability gate and D2 vacuously. For D3, its only free parameter is the emphasis baseline β, and sweeping it over B reverses 0.43–0.72% of pairs, so it passes on every dataset (`d3_pair_reversals.py`). Its activity boundary, however, is not robust; it is fixed. Sweeping the gate x_0 itself over B reverses 8.5–22.5% of pairs, against 8.8–22.6% for the entropy's baseline, and the two sweeps produce the same per-compound instability on Klaeger (mean rank $\sigma = 7.2$ and 7.5; Figure 6, right). The candidate's D3 standing therefore comes entirely from declaring x_0 a fixed constant, and Table 3 marks it $\checkmark^{\dagger}$. D4 holds by the proposition in Methods, and also on every compound in all four datasets with one kinase added at τ^*, the hardest case (`candidate_validation.py`; Supporting Information, Table S4).

A usable measure must also be recoverable from realistic panels. With the panel-size subsampling analysis above (same grid, 50 random subpanels per size), the candidate reaches mean Spearman $\rho > 0.90$ with its full-panel ranking at $p^* = 110$ kinases on Klaeger and $p^* = 80$ of 172 on Metz (Figure 7). This matches the entropy on both datasets; on Klaeger it is well ahead of the Gini (290) and the ratio (343), and on Metz the S-score converges faster

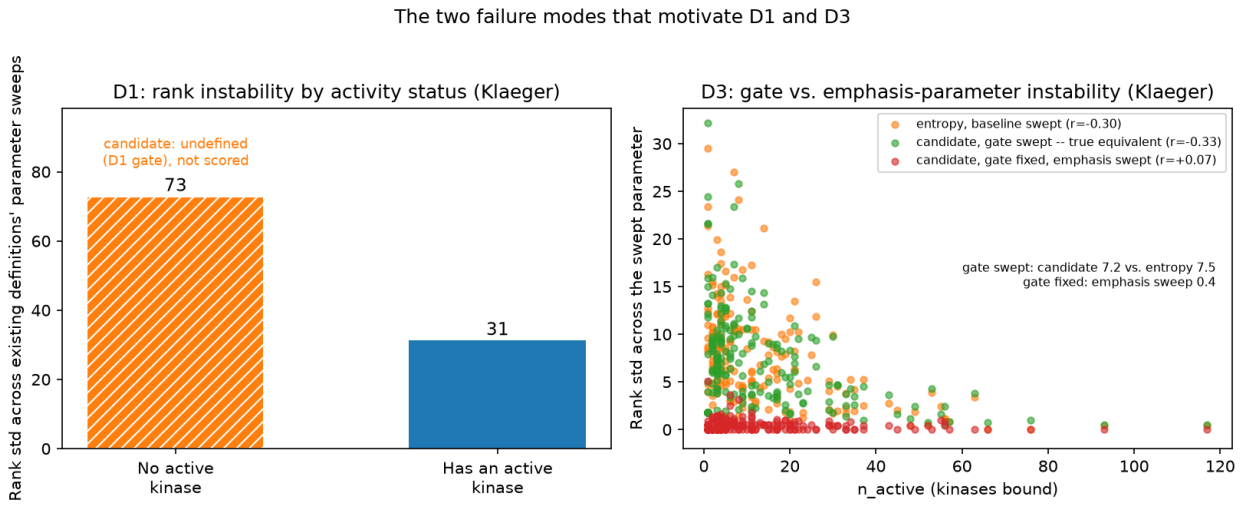


Figure 6: The two failure modes that motivate D1 and D3 (Klaeger). *Left (D1):* rank instability (std across the existing definitions' parameter sweeps) for compounds with no binding above the activity threshold versus compounds with at least one active kinase. The contrast is driven by the convention used for undefined entropy and Gini scores (Methods); computed only over the configurations in which each compound's score is defined, zero-active compounds are not less stable (text). *Right (D3):* per-compound rank instability against the number of active kinases, with each parameter swept over $B = [5.0, 6.0]$: the entropy's activity baseline (orange); the candidate's gate x_0, with its emphasis baseline pinned to x_0 at each point (green); and the candidate's emphasis baseline with x_0 fixed at the assay floor (red). The candidate's gate and the entropy's baseline produce nearly the same mean instability (7.2 and 7.5); only the emphasis sweep, with the gate held fixed, is stable (0.4).

(50). These p^* values score every compound, including ones the D1 gate would refuse to score once their active kinases fall outside a subpanel (42% of Klaeger compounds at 50 kinases, 27% at 110, at most 2% on Metz). Restricting the comparison to compounds that D1 allows to be scored on both the subpanel and the full panel moves the candidate, the entropy, and the Gini all to $p^* = 80$ on Klaeger, with Metz unchanged. Part of this gain is selection, since the gate removes the least-bound compounds, which are also the hardest to rank. The gate and the smooth emphasis thus neither cost nor buy convergence speed relative to the entropy.

The featured candidate is one member of a family of measures that satisfy D1–D4 on these data: a per-compound quantile floor in place of the fixed floor and Rényi orders up to $\alpha = 1$ do so on Davis and Klaeger, and smoothing widths from 0.75 to 2 times $\tau^* - x_0$ on Klaeger (D4 is proved only for the Shannon order, $\alpha = 1$). These variants, including how the

Rényi order affects panel recoverability, are reported in Supporting Information (Section S5).

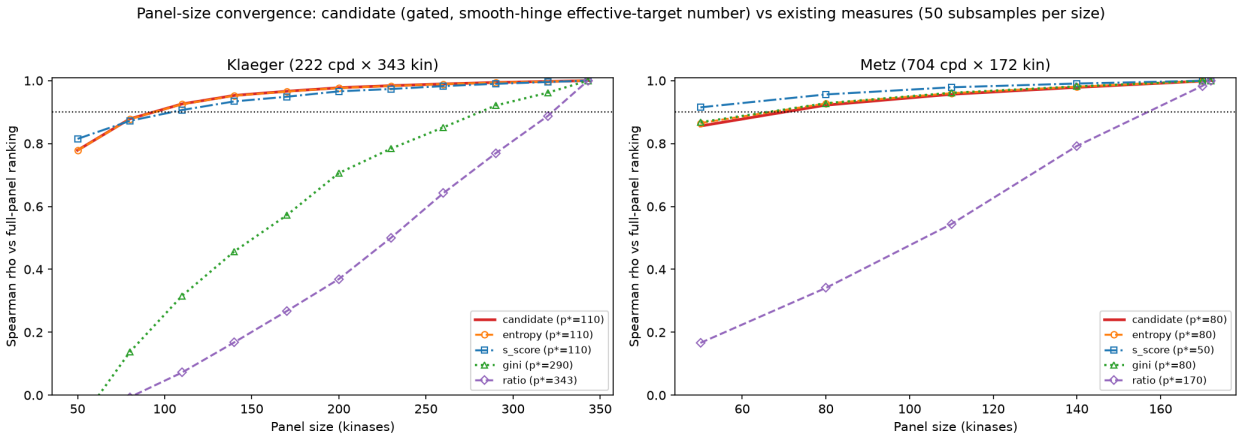


Figure 7: Panel-size convergence of the candidate measure compared with the four existing definitions, on the two sparse-coverage datasets (Klaeger, left; Metz, right; 50 random subpanels per size). Curves show the mean Spearman correlation between the ranking computed on a random subpanel and the full-panel ranking; p^* (legend) is the smallest panel reaching mean $\rho > 0.90$.

Why the Smooth Hinge Matters

The purpose of this subsection is to show that the candidate and the entropy differ in principle, although they rank these four datasets almost identically, and that the difference comes from the candidate's functional form rather than from a reporting rule (D1) or a declared constant (D3). The property at stake is D4.

The entropy can violate D4. Its hard cutoff gives a kinase just above the baseline β an arbitrarily small weight. A profile with many such kinases therefore has high entropy, and adding one kinase that is bound more strongly, but still below τ^*, concentrates the distribution and raises the selectivity score. For example, with $\beta = 5.0$ and $\tau^* = 6.0$, a profile with one kinase at $\mathrm{p}K_d = 7.0$ and fifty at 5.1 has an entropy-based selectivity of -4.89 bits, and adding one kinase at $\mathrm{p}K_d = 6.0$ raises it to -4.83. On Davis, Klaeger, and Metz, adding one kinase at τ^* never raises the entropy's score for any compound at any baseline in their sweeps. On Anastassiadis it does at low baselines: weighting each kinase by its raw percent inhibition ($\beta = 0\%$), as the published percent-inhibition forms do, it raises the score for 34%

of the 134 compounds with an active kinase, and for 16% at $\beta = 5\%$, although for none at 15% or above (`d4_entropy_vs_candidate.py`).

The candidate cannot. By the proposition in Methods, with $T \geq 0.58\,(\tau^* - x_0)$ no sub-threshold addition can raise its score. The condition on T is sharp: just below it, a profile with one strong target and enough kinases just above the floor violates D4. With $T = 1$ on Metz or $T = 10$ on Anastassiadis, both below the bound, constructed profiles of this kind violate D4 (by up to 0.014 bits on Anastassiadis), although none of the real compounds does. Setting $T = \tau^* - x_0$ on every dataset satisfies the condition with room to spare, leaves Davis and Klaeger unchanged, and changes the Metz and Anastassiadis rankings negligibly (Spearman 0.999 against $T = 1$ and $T = 10$, respectively).

The two measures rank these datasets almost identically, yet they differ on a question the rankings do not test: whether binding one more kinase below threshold can make a compound look more selective. For the entropy it can, for a third of the Anastassiadis compounds at the published forms' own baseline; for the candidate it provably cannot.

Discussion

The central finding of this work is that existing kinase inhibitor selectivity definitions, placed on a common log-affinity scale, do not measure the same quantity. The ratio definition and the distribution-based family (S-score, entropy, Gini) produce substantially different compound rankings (cross-family $r = 0.14$–0.62 across four datasets, versus within-family $r > 0.72$), and the sources of disagreement are mechanistically interpretable rather than random. This has practical consequences: a compound that ratio-based analysis designates as selective may be designated as moderately promiscuous by entropy-based analysis, and vice versa, depending on the shape of its binding profile. Go/no-go decisions in early drug discovery are routinely made on the basis of selectivity assessments without a principled choice of selectivity definition.

The clustering of definitions into two families reflects a difference in what is being measured. Distribution-based definitions measure the concentration of the compound's binding activity across the full profiled kinome. The ratio definition measures how much more potently the compound binds its primary target than a specified off-target. Among compounds that clear the activity threshold, the selectivity measure discrepancy is most pronounced for those with one clearly dominant target that also bind many kinases at moderate affinity (Type 3 profiles in our taxonomy, e.g. Dovitinib and BMS-690514). Ratio-based methods report these compounds as highly selective, because of the large gap to the next-best target. Distribution-based methods instead report them as moderately promiscuous, because activity is spread across many targets. Zero-active (Type 1) compounds have no target captured above the activity threshold on the profiled panel, so no definition can score them on measured binding, and D1 excludes them from scoring entirely. The properties D1–D4 (Methods) specify what we want any selectivity measure to do, with each property drawn from one of the failure modes above.

This coexistence has a historical explanation. In classical pharmacology and analytical chemistry, selectivity is inherently pairwise. IUPAC formalizes it as a selectivity coefficient, the ratio of binding constants for an analyte and an interferent,[26] and ratio measures entered kinase profiling on that basis. Kinome-wide profiling, however, poses a distributional question: how concentrated is a compound's binding across hundreds of kinases? Entropy and Gini were introduced because ratio measures are ill-suited to it, since they discard all but two points of a 300–500 kinase profile.

Practical Recommendations

Pending a formal resolution, three practical recommendations follow from our findings. First, D1 (reliability threshold) should be adopted as a reporting standard. Compounds with no binding above the $pK_d > 6$ activity threshold should not receive selectivity scores, or should have their scores explicitly flagged as unreliable. Of the 222 Klaeger compounds, 16 (7.2%) fall

into this category and currently receive scores whose ranks are set by an arbitrary convention for undefined values rather than by measured binding. Note that these compounds are not below the assay detection floor of $pK_d = 5.0$, which none of the 222 is: their maximum affinities run from 5.16 to 6.00. The gate is set at the activity threshold, not at the detection limit.

Second, selectivity assessments should report the sensitivity of rankings to parameter choice rather than a single score at a single parameter value. Our analysis shows that rank standard deviation across the parameter sweep is a compound-specific, interpretable quantity that identifies which compounds have robust selectivity assessments and which do not. This is analogous to reporting confidence intervals in statistical estimation.

Third, when comparing selectivity across studies that use different definitions, the comparison should be restricted to the subset of compounds and parameter values for which the definitions agree. Within the distribution-based family, the S-score, entropy, and Gini agree closely on the full panel ($r \geq 0.72$, the lowest being Klaeger S-score vs. Gini; Table 1; an observed value, not a preset bar, and distinct from the $\rho > 0.90$ panel-size criterion, which compares a definition with itself on a subpanel). They can be treated as interchangeable when parameters are chosen consistently and panels have similar, near-complete coverage. On smaller panels the Gini diverges: on Klaeger it needs 290 of 343 kinases to recover its full-panel ranking, against 110 for the entropy and the S-score. Much of that gap comes from scoring compounds that D1 would not score. Restricted to compounds that clear the activity threshold on both panels, the Gini recovers at 80 kinases, as does the entropy, though part of this gain is selection of the easier-to-rank compounds. Ratio-based measures, in turn, should not be compared directly with distribution-based measures at any panel size without acknowledging the conceptual distinction.

Limitations

The analysis covers four profiling datasets spanning three assay technologies and two readout types, but these still represent a fraction of the chemical and kinase space relevant to drug discovery. The assay technologies are not directly comparable in absolute terms — the active-compound fractions differ across datasets. As a consequence, our analysis relies on within-dataset rank comparisons rather than cross-dataset absolute scores. The two-family structure (in our common log-affinity operationalization) replicates across all four datasets, and the instability predictors across three of the four, which indicates that the findings are not artifacts of a single assay platform. The exception is the top_1–top_2 dependence on Anastassiadis ($r = -0.104$, $p = 0.23$), the one percent-inhibition dataset, so we cannot exclude that this predictor is specific to affinity readouts. However, findings should still be interpreted in light of these technological differences. Our comparison is also restricted to four widely used, historically established definitions. Bosc et al.'s window score and ranking score[8] were not included, but a check on the same four datasets confirms both correlate with the ratio family (Spearman $r = 0.80$–0.96) rather than forming a distinct third cluster. The CATDS score,[7] historically classified alongside the ratio-based measures, was also not implemented directly. Unlike $R(k)$, CATDS normalizes each target's binding reduction by the sum of reductions across the entire panel rather than comparing to a single off-target, a construction it shares with the distribution-based family rather than with $R(k)$. A proxy computed from our baseline-subtracted affinity profiles, in place of the raw concentration-response data CATDS requires, correlates far more strongly with Gini and entropy than with $R(k)$ on all four datasets (Spearman correlation 0.91–0.99 against Gini and 0.85–0.96 against entropy, versus 0.45–0.70 against $R(k)$), consistent with this structural reading. On our data CATDS does not behave like a ratio measure, and assigning it to that family would have misdescribed the cross-family correlations reported above. We did not add it as a fifth measure: the distribution family is already represented by three members, and CATDS appears to be an additional instance of the same underlying construct (full comparison numbers in Supporting

Information, Table S2, and in the code repository).

A further axis of platform dependence that this analysis does not examine is cell-based versus cell-free profiling. Recent work shows that selectivity classifications for the same compounds can differ substantially between biochemical binding panels and cellular target-engagement assays. For example, several type II (DFG-out-binding) inhibitors show cellular interactions that cell-free assays miss.[27,28] A plausible contributor, not established by that work, is that the cellular assays report occupancy of full-length kinase in intact cells, with regulatory domains, binding partners, and intracellular ATP present. Many cell-free binding panels, instead, score isolated, recombinantly expressed kinase-domain constructs. This specific contrast does not map cleanly onto the two cell-free datasets used here: Davis follows the isolated-domain convention, while Klaeger's Kinobeads approach instead captures endogenous, full-length kinases directly from cell lysate.

This raises the following question: for compounds profiled on more than one platform, does the choice of assay technology move a compound's selectivity ranking more or less than the choice of selectivity definition does? We resolved 46 compounds profiled in both Davis (competition binding) and Klaeger (chemoproteomics) by matching PubChem CIDs against Klaeger's drug names (`davis_klaeger_overlap.py`), and computed each definition's median rank on this fixed set of 46 compounds (separately on each platform). Rankings agree less across platforms than across definitions, so the *platform effect* (same definition, Davis rank vs. Klaeger rank) is on average the larger of the two: Spearman $r = -0.06$ to 0.68 across the four definitions for the platform comparison (mean 0.44), against $r = 0.14$ to 1.00 for the within-platform, cross-definition comparison (mean 0.61). The ratio definition is the extreme case: its Davis ranking and its Klaeger ranking for the same 46 compounds are statistically indistinguishable from uncorrelated ($r = -0.06$, $p = 0.67$). That is, which assay a compound happened to be profiled on matters at least as much as which selectivity definition is applied to it. This should temper any cross-study comparison of selectivity rankings regardless of which definition is used.

All four existing definitions and the candidate measure introduced above are symmetric functions of the binding profile: each treats a compound's affinity vector as an unordered multiset. Thus, two compounds with identical affinity distributions receive identical scores regardless of which specific kinases carry the off-target activity. A compound whose off-targets are pharmacologically inert and one whose off-targets are known safety-relevant kinases (e.g. a kinase implicated in a class toxicity) are therefore scored identically if their affinity distributions match. None of the measures discussed here incorporates target identity, and none accounts for exposure. A measured affinity says nothing about whether a target is actually occupied in a patient, since that depends on the drug concentration achievable at a tolerated dose, not on the pK_d or pK_i value alone. A selectivity score is consequently, at best, a summary of *binding breadth*. The go/no-go decisions should necessarily weigh it alongside target-specific liability knowledge and exposure modeling.

We also examined whether in vitro selectivity predicts clinical toxicity, correlating the selectivity rankings of the FDA-approved Klaeger compounds with FAERS adverse-event rates ($n = 46$) and FDA-label discontinuation rates ($n = 33$). None of the rate-based outcomes were significantly associated with selectivity under any definition (all $|r| < 0.17$, all $p > 0.27$). The only nominal signal is a weak correlation between raw FAERS report counts and the ratio ranking ($r = 0.31$, $p = 0.035$), which most likely reflects prescription volume rather than toxicity. We do not read the absence of a rate-based association as evidence that selectivity and toxicity are unrelated. FAERS report counts mostly track how often a drug is prescribed, and label discontinuation depends on the patient population treated for each indication, so neither can isolate the effect of selectivity. A proper test would require indication-stratified, treatment-related adverse-event rates from controlled trials with matched comparators, which are not available in standardized form across these drugs.

The candidate's distributional consistency rests on fixing its activity boundary at the assay floor x_0, and in none of the four datasets is x_0 a measured detection limit. On Davis and Klaeger it is the value assigned to pairs with no detected binding (p$K_d = 5.0$), on Metz the

lowest value reported in the released data ($pK_i = 4.0$), and on Anastassiadis a 20% inhibition noise level that we set, above the 0% to which missing and negative values are clipped. Fixing x_0 therefore removes the analyst's choice only to the extent that each dataset's encoding of non-detection is accepted as given. Where that encoding is itself uncertain, the candidate's standing on D3 is no better than the entropy's.

Finally, apart from the candidate's D4, which we prove, a formal proof of whether each definition satisfies each of D1–D4 (Table 3 gives the assessments) is left for future work. So is the question of whether the properties are jointly satisfiable only by measures within the candidate's family (Supporting Information, Section S5), and whether a uniqueness result analogous to Shannon's holds. A complementary direction is a consensus ranking that does not commit to any single measure, that is, a partial order on which all existing measures agree.

The four properties are also unlikely to be the complete set, and this work suggests three further candidates that we do not assess here. A measure should be correctly oriented against a crude external proxy for the concept, such as the number of active kinases, since a measure can satisfy D2 to D4 and still rank promiscuous compounds as selective. A measure should declare its scale type, because the four definitions report on incomparable scales and Figure 1 shows that leaving the operationalization undeclared can reverse a ranking. The panel-size behaviour reported above could itself be stated as a property. A further candidate, stability under measurement error in the affinities, is not testable on the present data. Agreement across datasets profiling the same compounds *is* testable, and is reported above as the platform-vs-definition comparison on the 46 compounds shared between Davis and Klaeger (by PubChem CID). By compound name, Klaeger and Anastassiadis share 19 compounds. However, a platform comparison between an affinity assay and a single-concentration percent-inhibition assay is less directly interpretable than the pKd-to-pKd Davis/Klaeger comparison above. Metz identifies compounds only by internal ID and PubChem SID rather than by name, and we did not resolve those against the other three datasets, so its overlap with Davis,

Klaeger, or Anastassiadis remains unquantified here.

Conclusion

We have shown that commonly used kinase inhibitor selectivity definitions, computed on a common log-affinity scale, split three to one, into a distribution-based group whose members agree closely and a ratio measure that agrees with none of them, and that the two measure fundamentally different properties of binding profiles. This grouping is a property of that common scale: under the association-constant form of the original publication, the entropy does not cluster with the S-score and Gini. We also showed that rank instability within and between families is associated with binding profile shape, and that three structurally distinct sources of instability can be identified and characterized. These findings motivate four properties that any well-formed selectivity measure should satisfy: a reliability threshold condition, bounded gap sensitivity, distributional consistency, and monotonicity under weak off-target addition.

None of the four definitions assessed satisfies all four properties. The one property all four violate is the reliability condition, a reporting convention that any of them could adopt. The candidate we propose, a gated, smooth-hinge effective-target number, satisfies all four: the reliability condition by a reporting rule; distributional consistency because its activity boundary is fixed at the assay floor (swept as a free parameter, that boundary moves the ranking as much as the entropy's baseline does); and monotonicity under weak off-target addition by its functional form, which we prove, and which the entropy fails for a third of the Anastassiadis compounds. The candidate recovers its full-panel ranking from roughly a third of a typical panel on Klaeger and roughly a half on Metz, as the entropy does, and it is part of a whole family of candidate definitions satisfying the four properties. Because the candidate is a local modification of the entropy, the two rank the compounds they both score almost identically (Spearman +0.999 on Klaeger), so our candidate measure is neither unique

nor an improvement on the entropy on these data. The case for the properties therefore does not rest on these data: they are requirements on every binding profile, and we propose them as checks that any new selectivity measure should pass.

The main goal of this work is to identify criteria for judging a candidate that could be considered a principled, universally adopted choice of selectivity measure in the future. Such criteria must be justified (which we argue they are), and jointly satisfiable (which the candidate demonstrates). Further properties may discriminate where these four do not, and the candidate may separate from the entropy on panels whose active fractions differ from the four studied here. Establishing additional properties that uniquely characterize a selectivity measure, as Shannon's axioms do for entropy, is a natural next step, for which additional datasets may be required.

The immediate practical implication is that selectivity assessments should report the sensitivity of rankings to parameter choice, and should explicitly flag compounds with no binding above the activity threshold as unreliable. The longer-term implication is that the field needs a principled basis for choosing among selectivity definitions. The need for a principled measure is not specific to selectivity. As future and ongoing work, we are developing a general measurement-theoretic framework for defining and comparing the competing measures of one concept, reporting where they agree and which pairs they leave undecided, with kinase selectivity as one instance alongside cases as varied as biological age and variant effect prediction.

Acknowledgement

The author thanks the developers of the Davis, Klaeger, Anastassiadis, and Metz datasets for making their data publicly available. The author used Claude (Anthropic, claude.ai), a large language model assistant, during the preparation of this manuscript. Specifically, Claude was used to assist with: data analysis code development and debugging in Python;

literature search and identification of related work; drafting and editing of manuscript text; and research planning and interpretation of results. All scientific conclusions, the conceptual framework, the motivated properties D1–D4, and the final manuscript content were developed, verified, and approved by the author. The author takes full responsibility for the accuracy and integrity of all content in this work. The use of Claude does not transfer any rights over the content of this manuscript to Anthropic or any third party.

Data and Software Availability

All data and software used in this study are publicly available. The Davis kinase inhibitor profiling dataset is available at `https://github.com/dingyan20/Davis-Dataset-for-DTA-Prediction`. The Klaeger et al. dataset is available as supplementary data to the original publication (Science, 2017, `https://doi.org/10.1126/science.aan4368`). The Anastassiadis et al. (*Nat. Biotechnol.* 2011) and Metz et al. (*Nat. Chem. Biol.* 2011) datasets are available as supplementary data to their original publications (`https://doi.org/10.1038/nbt.2017`; `https://doi.org/10.1038/nchembio.530`). All analysis scripts, processed data files, and the LaTeX source for this manuscript are available at `https://github.com/polinavino/kinase-selectivity-definitions`; this repository includes the code and full per-definition results for the secondary clinical-outcome analysis (FAERS adverse-event report rates, obtained via openFDA queries, and FDA-label discontinuation rates, from FDA product labels, for the FDA-approved Klaeger compounds). No confidential data were used in this study. All results reported in this manuscript are fully reproducible from the publicly available data and code. A preprint of this manuscript is available on ChemRxiv at `https://doi.org/10.26434/chemrxiv.15001618/v1`.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Conflict of Interest

The author declares no conflict of interest.

References

(1) Roskoski, R. Properties of FDA-approved small molecule protein kinase inhibitors: a 2024 update. *Pharmacol. Res.* **2024**, *200*, 107059.

(2) Manning, G.; Whyte, D. B.; Martinez, R.; Hunter, T.; Sudarsanam, S. The protein kinase complement of the human genome. *Science* **2002**, *298*, 1912–1934.

(3) Bhullar, K. S.; Lagarón, N. O.; McGowan, E. M.; Parmar, I.; Jha, A.; Hubbard, B. P.; Rupasinghe, H. P. V. Kinase-targeted cancer therapies: progress, challenges and future directions. *Mol. Cancer* **2018**, *17*, 48.

(4) Karaman, M. W.; Herrgard, S.; Treiber, D. K.; et al. A quantitative analysis of kinase inhibitor selectivity. *Nat. Biotechnol.* **2008**, *26*, 127–132.

(5) Uitdehaag, J. C. M.; Zaman, G. J. R. A theoretical entropy score as a single value to express inhibitor selectivity. *BMC Bioinformatics* **2011**, *12*, 94.

(6) Graczyk, P. P. Gini coefficient: a new way to express selectivity of kinase inhibitors against a family of kinases. *J. Med. Chem.* **2007**, *50*, 5773–5779.

(7) Klaeger, S.; Heinzlmeir, S.; Wilhelm, M.; et al. The target landscape of clinical kinase drugs. *Science* **2017**, *358*, eaan4368.

(8) Bosc, N.; Meyer, C.; Bonnet, P. The use of novel selectivity metrics in kinase research. *BMC Bioinformatics* **2017**, *18*, 17.

(9) Bello, T.; Gujral, T. S. KInhibition: a kinase inhibitor selection portal. *iScience* **2018**, *8*, 49–53.

(10) Shannon, C. E. A mathematical theory of communication. *Bell Syst. Tech. J.* **1948**, *27*, 379–423.

(11) Miljkovic, F.; Bajorath, J. Reconciling selectivity trends from a comprehensive kinase inhibitor profiling campaign with known activity data. *ACS Omega* **2018**, *3*, 17196–17206.

(12) Saifudeen, M.; Zhu, S.; Liang, S.; Eason, M.; Goupil, A.; Mische, D. F.; Loch, C. M.; Ma, H.; Chan, M.; Gujral, T. S. Comprehensive profiling of clinically approved kinase inhibitors reveals mutation-specific inhibitors and opportunities for drug repurposing. *Nat. Biotechnol.* **2026**,

(13) Khinchin, A. Y. *Mathematical Foundations of Information Theory*; Dover: New York, 1957.

(14) Arrow, K. J. *Social Choice and Individual Values*; Wiley: New York, 1951.

(15) Singleton, J.; Booth, R. Towards an axiomatic approach to truth discovery. *Auton. Agents Multi-Agent Syst.* **2022**, *36*, 43.

(16) Davis, M. I.; Hunt, J. P.; Herrgard, S.; et al. Comprehensive analysis of kinase inhibitor selectivity. *Nat. Biotechnol.* **2011**, *29*, 1046–1051.

(17) Öztürk, H.; Özgür, A.; Ozkirimli, E. DeepDTA: deep drug–target binding affinity prediction. *Bioinformatics* **2018**, *34*, i821–i829.

(18) Anastassiadis, T.; Deacon, S. W.; Devarajan, K.; Ma, H.; Peterson, J. R. Comprehensive assay of kinase catalytic activity reveals features of kinase inhibitor selectivity. *Nat. Biotechnol.* **2011**, *29*, 1039–1045.

(19) Metz, J. T.; Johnson, E. F.; Soni, N. B.; Merta, P. J.; Kifle, L.; Hajduk, P. J. Navigating the kinome. *Nat. Chem. Biol.* **2011**, *7*, 200–202.

(20) Gao, Y.; Davies, S. P.; Augustin, M.; Woodward, A.; Patel, U. A.; Kovelman, R.; Harvey, K. J. A broad activity screen in support of a chemogenomic map for kinase signalling research and drug discovery. *Biochem. J.* **2013**, *451*, 313–328.

(21) Patricelli, M. P. et al. In situ kinase profiling reveals functionally relevant properties of native kinases. *Chem. Biol.* **2011**, *18*, 699–710.

(22) Harris, C. R.; Millman, K. J.; van der Walt, S. J.; et al. Array programming with NumPy. *Nature* **2020**, *585*, 357–362.

(23) McKinney, W. Data structures for statistical computing in Python. Proceedings of the 9th Python in Science Conference. 2010; pp 56–61.

(24) Virtanen, P.; Gommers, R.; Oliphant, T. E.; et al. SciPy 1.0: fundamental algorithms for scientific computing in Python. *Nat. Methods* **2020**, *17*, 261–272.

(25) Koo, T. K.; Li, M. Y. A guideline of selecting and reporting intraclass correlation coefficients for reliability research. *J. Chiropr. Med.* **2016**, *15*, 155–163.

(26) Vessman, J.; Stefan, R. I.; van Staden, J. F.; et al. Selectivity in analytical chemistry. *Pure Appl. Chem.* **2001**, *73*, 1381–1386.

(27) Binder, M. J.; Bashore, F. M.; Dunn Hoffman, K. K.; Damgaard, C.; Slater, M.; Drewry, D. H.; Robers, M. B.; Axtman, A. D. Cellular context influences kinase inhibitor selectivity. *J. Med. Chem.* **2026**, *69*, 4208–4220.

(28) Binder, M. J.; Bashore, F. M.; Dunn Hoffman, K. K.; Damgaard, C.; Slater, M.; Drewry, D. H.; Robers, M. B.; Axtman, A. D. Correction to "Cellular context influences kinase inhibitor selectivity". *J. Med. Chem.* **2026**, *69*, 20069–20071.