Supporting Information

Analysis Choices and Ranking Stability in Kinase Inhibitor Selectivity Scores

Polina Vinogradova

This document gives the full evidence for results summarized in the main text, in the order of its argument: dataset construction (Section S1), score implementations and agreement (S2), ranking sensitivity and profile associations (S3), monotonicity under weak-target addition (D3) and the gated entropy (S4), recovery from subpanels (S5), cross-dataset and supplementary comparisons (S6), the clinical analysis (S7), and the script behind every result (S8). All analyses use the curated panels of the main text, one reference target per human protein kinase, except where Section S1 compares them with the released panels. The code repository (`https://github.com/polinavino/kinase-selectivity-definitions`) reproduces every number, and the script behind each table and paragraph is named next to it.

S1: Dataset construction and target selection

Released and curated panels. Table S1 compares the main Davis and Klaeger results on the released and curated panels. The released Davis panel lists 433 constructs, 12 of them for EGFR. For 26 of the 68 compounds the two best-bound constructs are forms of one gene. The most ratio-unstable compound on the released panel, GDC-0879, owes its 0.023 top_1–top_2 gap to BRAF(V600E) and BRAF, and its gap to the next gene (RAF1) is 1.54. On Davis, the best affinity per gene is also given as an alternative to the reference construct.

Klaeger et al. annotate 225 of the 343 released Klaeger targets as protein or lipid kinases. The curated panel keeps 217 of them and leaves out three lipid-kinase columns (PIP4K2A, PIP4K2C and PIP4K2A;PIP4K2B), four protein groups (CSNK2A1;CSNK2A3, PDPK1;PDPK2P, PRKCE; PRKCH and PRKX;PRKY), and Q6ZSR9, which UniProt describes as an uncharacterized protein with no kinase domain, encoded in the 3'-UTR of AAK1. The other 118 targets are not annotated as kinases: 46 subunits or adaptors of kinase complexes (for example the cyclins CCNA2–CCNT2, CKS1B, the AMPK and PKA regulatory subunits, and the AP2 adaptor subunits), 7 small-molecule kinases, 20 other enzymes (for example COMT, NQO2, TOP2B, PYGB, PYGL and PTPN1), 43 other proteins, including ribosomal proteins and SLC transporters, and BCR and TRRAP. The source flags 88 of the 118 as possibly co-enriched as complex members, among them BCR (with ABL1) and TRRAP (with CDK8). The 122 targets outside the curated panel, other than the four protein groups, hold 556 of the 2,763 active entries (20.1%), in 122 of the 222 compounds. Six compounds are active only on them (AZD-2014, AZD-8186, NMS-1286937, SGX-523, TG-100115 and volitinib). The code repository lists every target of each dataset with the reason it is kept or removed. Most results change little between the panels. The largest changes are on Davis, where the ratio agrees more with the other definitions on the curated panel, and on Klaeger, where with all compounds included the Gini correlates more with the ratio than with the S-score on the curated panel. The p^* rows of Table S1 are from the analysis without the D1 gate (Section S5), not from the main-text analysis.

Table S1: Main Davis and Klaeger results on the released and curated panels. Correlations are Spearman r; "active" rows re-rank the compounds with an active kinase among themselves, as do the Table S4 rows. p^* is from the panel-size analysis without the D1 gate (Section S5). (`davis_gene_level.py`, `klaeger_protein_kinases.py`)

Davis	Released (433)	Curated (363)	Best per gene (363)
Same-gene top two constructs	26	0	0
Table 3: S–E, S–G, E–G	0.946, 0.952, 0.999	0.963, 0.963, 0.999	0.960, 0.960, 1.000
Table 3: R against S, E, G	0.456, 0.343, 0.343	0.563, 0.513, 0.510	0.511, 0.478, 0.474
Table S4: n_{active}, entropy	-0.279	-0.301	-0.268
Table S4: n_{active}, ratio	$+0.294$	$+0.007$	$+0.189$
Table S4: top_1–top_2 gap, ratio	-0.444	-0.445	-0.566
Table S4: top_2 pK_d, ratio	$+0.296$	$+0.147$	$+0.203$
Gated entropy vs. entropy	0.999	0.997	0.998
p^* without D1 gate: gated E, E, S, G, R	50, 50, 50, 110, 410	50, 50, 50, 110, 350	50, 50, 50, 110, 350

Klaeger	Released (343)	Curated (217)
Zero-active compounds	16	22
Table 3, all: S–E, S–G, E–G	0.910, 0.722, 0.840	0.896, 0.508, 0.658
Table 3, all: R against S, E, G	0.223, 0.455, 0.616	0.078, 0.344, 0.655
Table 3, active: S–E, S–G, E–G	0.913, 0.911, 0.998	0.897, 0.892, 0.997
Table 3, active: R against S, E, G	0.418, 0.621, 0.634	0.378, 0.622, 0.635
$r(\text{S},\text{G}) - r(\text{G},\text{R})$, active, 95% interval	$[0.195, 0.375]$	$[0.158, 0.355]$
Aggregated rank σ, zero-active vs. active	74.0 vs. 31.2	77.9 vs. 32.2
Table S4: n_{active}, entropy	-0.302	-0.312
Table S4: top_1–top_2 gap, ratio	-0.291	-0.275
Table S4: active range, S-score	-0.315	-0.317
Table S4: skewness, Gini	$+0.107$	$+0.072$
Gated entropy vs. entropy	0.998	0.998
p^* without D1 gate: gated E, E, S, G, R	80, 80, 110, 290, > 320	80, 80, 80, 200, > 200

Klaeger confidence filter. The high-confidence filter removes 793 of the 5,916 drug–target rows (13.4%), one of which has no numeric K_d^{app}. A removed pair is coded at the floor by convention; a low-confidence fit does not establish weak binding. Of the 792 low-confidence pairs with a numeric K_d^{app}, 405 involve curated targets (375 with a fitted pK_d above 5, and 208 above 6), and 402 of these fall within the 222×217 matrix. Adding them changes 0.77% of its entries (some added values lie below the floor and are set to it). It shifts the all-compound Klaeger correlations of main-text Table 3 by at most 0.08, and those among the 200 compounds with an active kinase under both filters by at most 0.006. It leaves every p^* of the main-text recovery analysis unchanged and raises coverage at 50 targets from 70.6% to 72.9%. Retaining all low-confidence pairs also adds 7 compounds and shifts the all-compound correlations by at most 0.06. (`klaeger_filter_sensitivity.py`)

Metz coverage. The 704 retained Metz compounds are measured against 50 to 143 of the 172 kinases (29.1% to 83.1%; mean 40.9%, median 37.2%, 75th percentile 47.1%). For the median compound, most entries are therefore floor-imputed. Twenty-three compounds have exactly 50 measured kinases and 170 have 55 or fewer. Only 7 kinases are measured for at least 80% of the compounds, so no fully measured subset of useful size exists. (`metz_coverage.py`, `claim_checks.py` section 10)

Activity threshold on Davis. Active entries ($\mathrm{p}K_d > 6$) make up 16.7% of the Davis matrix and 4.5% of the Klaeger matrix; the 46 compounds profiled in both are active against 15.0% of the Davis panel and 3.9% of the Klaeger panel on average. With the S-score's threshold sweep shifted up by one log unit on Davis (from $[5.5, 8.0]$ to $[6.5, 9.0]$), the S-score still correlates more with the entropy and Gini (0.88) than with the ratio (0.54). (`dataset_descriptive_stats.py`, `claim_checks.py` section 9)

Datasets not used. The screen of Gao et al. (2013) has the same readout type as Anastassiadis, so it would add no new assay technology, and its raw data are not openly redistributable. The study of Patricelli et al. (2011) profiles too few inhibitors for a ranking analysis.

S2: Score implementations and agreement

Table S2 gives the correlation of the published, association-constant entropy with the median rankings of the S-score, the log-affinity Gini and the ratio, with floor-coded pairs given zero weight and kept at the floor, and the native percent-inhibition Gini on Anastassiadis. (`original_operationalization_robustness.py`)

Table S2: Spearman correlation of the published entropy (association-constant weighting) with the median rankings of the other definitions.

Dataset	Floor-coded pairs	S-score	Gini	Ratio
Davis ($n = 68$)	zero weight	0.738	0.736	0.861
	at the floor	0.344	0.368	0.688
Klaeger ($n = 222$)	zero weight	0.571	0.530	0.657
	at the floor	−0.597	0.077	0.648
Klaeger, active ($n = 200$)	zero weight	0.564	0.782	0.916
	at the floor	−0.412	−0.094	0.591
Metz ($n = 704$)	zero weight	0.291	0.385	0.758
	at the floor	0.241	0.342	0.742

Anastassiadis ($n = 178$): native percent-inhibition Gini (no baseline subtraction) against the median rankings of the other definitions.

	S-score	Entropy	Ratio
Gini (native)	0.340	0.629	0.326

In every case the published entropy correlates more with the ratio than with the S-score or the Gini, so it would not be grouped with them on any of the three affinity datasets. How far it is from the S-score and Gini depends on the treatment of floor-coded pairs. On Klaeger, a third treatment keeps every measured high-confidence affinity at its measured value, including the 194 weaker than the floor (4,000 measured cells on the curated panel, in 84 compounds for the weak ones), and omits unmeasured pairs. The published entropy then correlates with the S-score, Gini and ratio at 0.558, 0.531 and 0.663 over all compounds and 0.558, 0.778 and 0.918 over the 200 with an active kinase, against 0.571, 0.530, 0.657 and 0.564, 0.782, 0.916 with floor-coded pairs given zero weight. The two rankings correlate at 0.9987 over all compounds and 0.9996 over active ones. On the subpanels of the recovery analysis, p^* stays at 200 targets, and the mean Spearman correlation at 50 targets is 0.6025 against 0.6014.

Figure S1 shows the correlations between the four definitions on all four datasets, among compounds with an active kinase and among all compounds. Among the S-score, entropy and Gini, the correlations among active compounds are 0.963–0.999 (Davis), 0.892–0.997 (Klaeger), 0.870–0.999 (Anastassiadis) and 0.806–0.994 (Metz), and the ratio's correlations with them 0.510–0.563, 0.378–0.635, 0.793–0.843 and 0.147–0.203. With all compounds, the ranges are 0.508–0.896 and 0.078–0.655 on Klaeger, 0.859–0.988 and 0.478–0.553 on Anastassiadis, and 0.810–0.994 and 0.138–0.193 on Metz. Using mean instead of median ranks changes no all-compound correlation by more than 0.036 (Davis), 0.076 (Klaeger), 0.075 (Anastassiadis) or 0.052 (Metz). (`additional_datasets_analysis.py`, `claim_checks.py` sections 3 and 11)

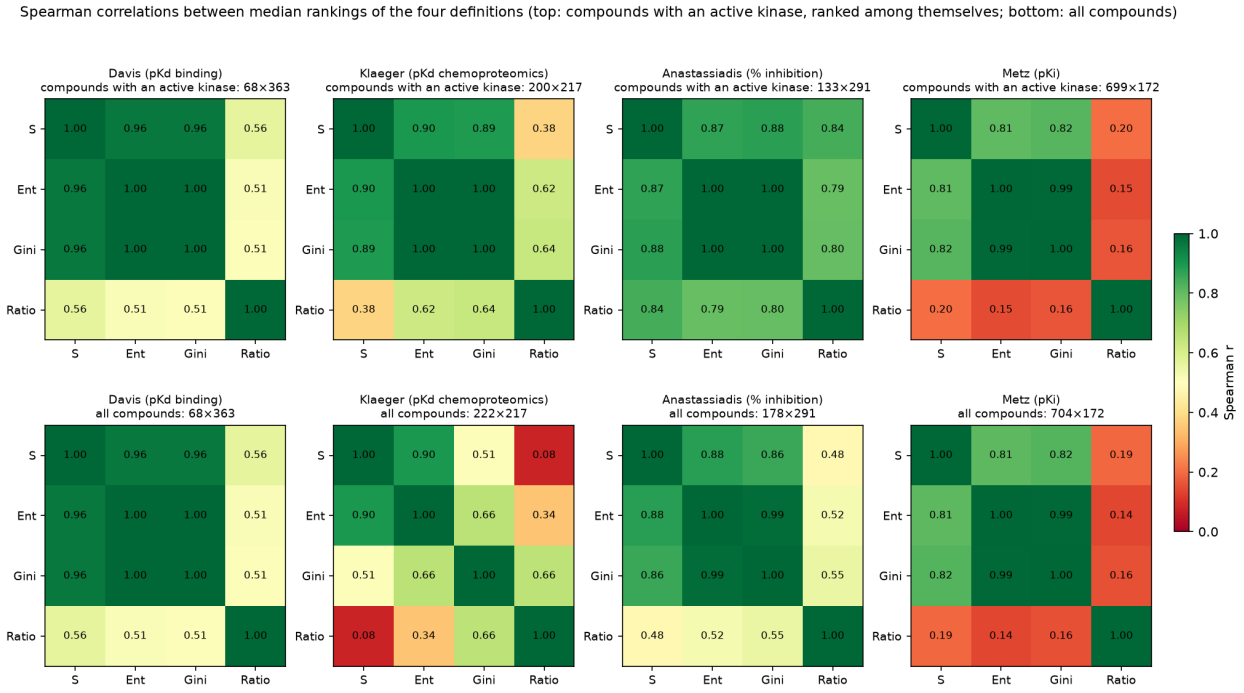


Figure S1: Spearman correlations between the median rankings of the four definitions on each dataset (compounds × targets above each panel). Top: compounds with an active kinase, ranked among themselves. Bottom: all compounds. S: S-score; Ent: entropy.

S3: Ranking sensitivity and profile associations

Compounds with no active kinase. Compounds with no active kinase have higher aggregated rank instability than active compounds on every dataset where they occur (77.9 against 32.2 on Klaeger, 36.6 against 21.8 on Anastassiadis, and 178.5 against 108.1 on Metz, which has five such compounds). The contrast comes mainly from the definitions disagreeing. Within a single definition's sweep on Klaeger, it is significant only for the Gini (65.0 against 16.3, $p < 0.001$). It is small and not significant for the entropy (11.9 against 12.6, $p = 0.91$) and the ratio (17.5 against 19.1, $p = 0.51$), and reverses for the S-score (19.0 against 25.7, $p = 0.01$), whose score is 0 for every such compound at thresholds of 6 or above. On Anastassiadis it is significant for the Gini, the entropy and the ratio. On Metz it reverses for the S-score, the entropy and the Gini (7.8 against 71.2, 11.0 against 24.5, and 5.4 against 24.2; all $p < 0.05$). On Klaeger, the aggregated contrast persists with the undefined entropy and Gini configurations dropped (66.2 against 31.7) and with only the S-score and ratio aggregated (70.6 against 37.7, both $p < 0.001$). Computed only over each compound's defined configurations, the 16 Klaeger compounds with no active kinase that have at least two defined configurations have lower mean instability than active ones within the entropy (7.0 against 13.5) and the Gini (6.7 against 13.6, both $p = 0.002$). This is a group result: three of the 16 (brivanib, brivanib alaninate and NMS-1286937) are less stable than the active mean, and the compounds contribute different numbers of configurations. (`claim_checks.py` sections 13 and 20, `undefined_score_convention.py`)

Six Klaeger compounds with no detected binding. AZD-2014, AZD-8055, AZD-8186, BMS-911543, SGX-523 and volitinib show no binding to any curated target, although they are established inhibitors of mTOR, mTOR, PI3Kβ, JAK2, MET and MET. mTOR and PI3Kβ were not quantified for any of the 229 compounds with a fitted target in the Kinobeads table, consistent with low expression or poor capture in the lysate; PI3Kβ, a lipid kinase, is also outside the curated panel. JAK2 and MET were quantified for 7 and 32 other compounds (up to p$K_d = 8.3$ and 9.5), so absence from the lysate does not explain the result for BMS-911543, SGX-523 and volitinib. These data cannot tell whether compound-specific capture, competition or curve fitting does. (`claim_checks.py` section 1)

Boundary sensitivity. Table S3 gives the share of pairs of scored compounds whose order reverses as each parameter moves across $B = [x_0, \tau^*]$. B is $[5.0, 6.0]$ on Davis and Klaeger, $[4.0, 6.0]$ on Metz and $[20, 50]$% inhibition on Anastassiadis. For the ratio, the parameter is the level b in $R_b(k) = x_{(1)} - \max(x_{(k+1)}, b)$. The gated entropy's gate is swept with its emphasis baseline moved with it and T held at its value for the assay floor. A score satisfies D2 only if it stays within $\varepsilon = 1\%$ on every dataset and, for the ratio, for every k. Every swept parameter except the gated entropy's emphasis baseline exceeds 13% on at least one dataset (the ratio's level only at $k = 5$). Figure S2 shows the per-compound instability behind the Klaeger values. (`d3_pair_reversals.py`, `failure_mode_illustration.py`)

Profile features and instability. Table S4 gives the correlations between profile features and within-definition instability, with the compounds that have an active kinase ranked among themselves, and, for Klaeger, ranked among all 222 compounds. The ranking population changes the Gini column most: ranked among all compounds, the skewness of the active affinities is the Gini's strongest correlate (+0.262), and ranked among active compounds it is not significant (+0.072), while the number of active kinases becomes significant (−0.254). Restricted to the 189 active

Table S3: D2: pairs of scored compounds reversing in B (%), for every swept parameter. The S-score is evaluated exactly. The other parameters are evaluated on 201 points, and the gate also at every observed value in B and just below it, where its score jumps; these values are grid estimates (101 points change no value by more than 0.1 percentage point).

Parameter swept	Davis	Klaeger	Anastassiadis	Metz
S-score, threshold τ	8.43	9.96	12.99	21.82
Entropy, baseline β	7.90	12.37	10.58	22.80
Gini, baseline β	8.56	12.40	10.51	22.49
Gated entropy, emphasis β (gate fixed)	0.44	0.53	0.36	0.72
Gated entropy, gate x_0	9.22	12.99	14.39	23.59
Ratio, level b, $k = 1$	0.00	3.00	2.13	0.04
Ratio, level b, $k = 2$	0.83	6.26	4.10	0.43
Ratio, level b, $k = 3$	1.10	9.19	6.08	0.75
Ratio, level b, $k = 4$	1.71	12.03	8.17	0.84
Ratio, level b, $k = 5$	2.59	13.71	9.72	1.03

compounds whose entropy and Gini are defined throughout the baseline sweep, and re-ranked, the number of active kinases correlates with entropy and Gini instability at -0.274 and -0.272 (both $p < 0.001$). Figure S3 plots each definition's strongest correlate. With a Bonferroni correction for the 24 tests per dataset, 8 of the 13 nominally significant Klaeger correlations survive, and on Davis only the ratio–gap correlation. The Holm procedure gives the same result; the Benjamini–Hochberg procedure keeps 13 Klaeger correlations and, on Davis, again only the ratio–gap correlation. On Anastassiadis and Metz, with active compounds ranked among themselves, the number of active kinases correlates with entropy instability at -0.38 and -0.42 (both $p < 0.001$), and the top_1–top_2 gap with ratio instability at 0.00 ($p = 0.97$) and -0.26 ($p < 0.001$). Without the eight Anastassiadis compounds whose two best targets are both at 99.5% inhibition or more, the gap–ratio correlation there is -0.15 ($p = 0.10$). (`instability_by_family.py`, `additional_datasets_analysis.py`, `claim_checks.py` sections 4 and 21)

Undefined features. The skewness of the active affinities is undefined with fewer than three active kinases, and their standard deviation and range with fewer than two; Table S4 sets these to 0. Over only the compounds for which each feature is defined, the Klaeger range associations strengthen (S-score, entropy and Gini -0.373, -0.378 and -0.349, all $p < 0.001$), the Gini–skewness association stays non-significant ($+0.145$, $p = 0.08$), and no Davis correlation changes significance. (`feature_placeholder_sensitivity.py`)

Range and skewness. Range rises with the number of active kinases (Spearman 0.76 on Klaeger). Holding that number fixed, range remains associated with S-score instability (partial $r = -0.22$, $p = 0.002$) but not with entropy or Gini instability ($+0.08$ and $+0.01$). On the 46 compounds profiled in both Davis and Klaeger, the range–S-score association is weak on both (Klaeger $r = -0.10$, $p = 0.56$, for the 37 with at least two active kinases; Davis -0.16, $p = 0.30$); the same compounds have a median of 4 active kinases on Klaeger and 32 on Davis. Among the 148 Klaeger compounds with at least three active kinases, skewness correlates with the share of a compound's weight carried by kinases that the β sweep crosses (0.38), and that share with Gini instability with the number of active kinases held fixed (partial $r = 0.43$). Ranked among active compounds, skewness has no marginal association with Gini instability, but with the number of active kinases held fixed it does (partial $r = 0.35$, $p < 0.001$); holding the weight share fixed as well lowers this to 0.19

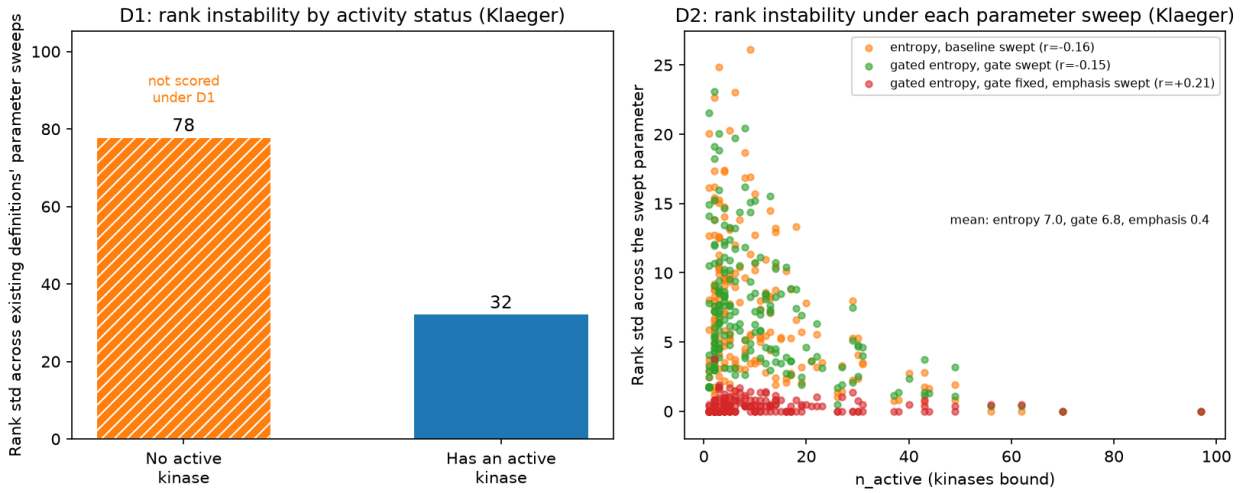


Figure S2: Rank instability on Klaeger. *Left (D1):* rank standard deviation across the four definitions' parameter sweeps, for the compounds with no binding above the activity threshold and for those with an active kinase, all ranked together. *Right (D2):* per-compound rank standard deviation, among compounds with an active kinase ranked among themselves, against the number of active kinases, with one parameter swept over $B = [5.0, 6.0]$: the entropy's baseline β, the gated entropy's gate x_0 with its emphasis baseline moved with it, and the gated entropy's emphasis baseline with x_0 fixed at the floor.

($p = 0.02$). The share of a compound's above-floor kinases that some baseline crosses correlates with its entropy instability ($r = 0.22$, $p = 0.002$). (`claim_checks.py` sections 8 and 15–17)

Further features. Five further features were tested on Klaeger and Davis: the top-target concentration ratio CR_1 and the Herfindahl–Hirschman Index (HHI) from industrial-organization economics, the primary-target affinity, a near-tie count generalizing the top_1–top_2 gap, and the active-affinity coefficient of variation. CR_1 and HHI are near-collinear with n_{active} ($r = -0.98$ for CR_1, and -0.997 and -0.998 for HHI). The other three were weaker than the strongest feature of Table S4 for every definition and on both datasets, except the primary-target affinity, which correlates with Gini instability on Klaeger at -0.295. (`instability_predictor_candidates.py`)

Unstable and stable compounds. All 20 most unstable Klaeger compounds, ranked together, have no active kinase. Among the 20 most unstable compounds with an active kinase, ranked among themselves, five have a top_1–top_2 gap below 0.1, five have at least 10 active kinases, and ten have neither. Of the ten largest entropy–ratio disagreements among active Klaeger compounds, seven have 14–27 active kinases, six of them with a top_1–top_2 gap above 1, for example dovitinib ($n_{\mathrm{active}} = 22$, gap 1.93) and BMS-690514 (20, 1.98). The ratio ranks these seven 31st to 63rd of 200, and the entropy and Gini 146th to 180th. The other three have one to three active kinases and are ranked as more selective by the entropy: fasudil (3, gap 0.03; ratio rank 181, entropy 68), CUDC-101 (3, 0.19; 172, 61) and rigosertib (1, 1.03; 118, 10). The most stable active Klaeger compounds are among the most promiscuous: XL-228 ($\sigma = 3.0$, 49 active kinases) and PF-03814735 (5.4, 40) have median ranks of 190–198 of 200 under every definition, and mean instability falls from 25–32 for compounds with 1 to 40 active kinases to 15 above 40. On Davis, compounds with few active kinases are the most stable on average (mean $\sigma = 3.9$ for one or two, 4.5 for three to five, and 7.3

Table S4: Spearman correlations between profile features and within-definition rank instability. Top: Klaeger (200 compounds with an active kinase, ranked among themselves) and Davis (68). Bottom: the same 200 Klaeger compounds, ranked among all 222. S: S-score; E: entropy; G: Gini; R: ratio. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

	Klaeger, active ranking				Davis			
Feature	S	E	G	R	S	E	G	R
n_{active}	-0.236^{***}	-0.312^{***}	-0.254^{***}	-0.040	-0.210	-0.301^{*}	-0.313^{**}	$+0.007$
top_1–top_2 gap	$+0.074$	$+0.019$	$+0.016$	-0.275^{***}	$+0.043$	$+0.074$	$+0.058$	-0.445^{***}
Active affinity std	-0.254^{***}	-0.083	-0.135	-0.129	-0.017	$+0.020$	$+0.035$	-0.176
Active affinity range	-0.317^{***}	-0.185^{**}	-0.188^{**}	-0.118	-0.023	-0.065	-0.076	-0.240^{*}
top_2 pK_d	-0.271^{***}	-0.157^{*}	-0.218^{**}	$+0.188^{**}$	-0.174	-0.186	-0.190	$+0.147$
Active affinity skew	-0.199^{**}	-0.051	$+0.072$	-0.089	$+0.041$	$+0.195$	$+0.176$	-0.113

	Klaeger, ranked among all 222			
Feature	S	E	G	R
n_{active}	-0.253^{***}	-0.394^{***}	-0.075	-0.025
top_1–top_2 gap	$+0.016$	$+0.057$	$+0.057$	-0.342^{***}
Active affinity std	-0.308^{***}	-0.132	-0.056	-0.121
Active affinity range	-0.370^{***}	-0.247^{***}	-0.054	-0.113
top_2 pK_d	-0.300^{***}	-0.230^{**}	-0.219^{**}	$+0.221^{**}$
Active affinity skew	-0.236^{***}	-0.096	$+0.262^{***}$	-0.107

for more than 40). The most stable, nintedanib ($\sigma = 2.1$), has 153 active targets and median ranks of 59–62 of 68. Without the ratio's configurations, the mean instability of the 32 Davis compounds with more than 40 active targets falls from 7.3 to 3.7. (`claim_checks.py` sections 5–7 and 18)

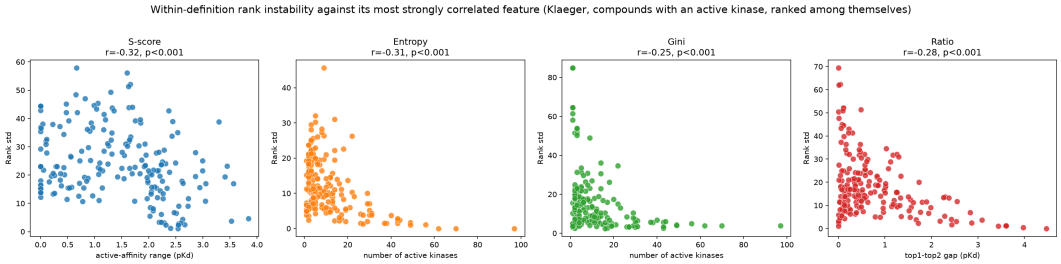


Figure S3: Within-definition rank instability against the profile feature most strongly correlated with it (Klaeger, the 200 compounds with an active kinase, ranked among themselves): active-affinity range (S-score), number of active kinases (entropy and Gini), and top_1–top_2 gap (ratio). Spearman r above each panel.

S4: Monotonicity under weak-target addition and the gated entropy

Proof of the proposition. Every gated kinase, the added one included, has weight above $m = T \ln 2$. Since the softplus increases with affinity, an added kinase with value at most τ^* has weight $a \le T \ln(1 + e^{(\tau^* - x_0)/T})$, and by the assumption on T this is at most $e\,m$. Let the original profile have $n \ge 1$ gated kinases ($n \ge 1$ because D1 scores it only if some kinase lies above $\tau^* > x_0$), total weight $W > nm$, and entropy H (in nats), and let $q = a/(W + a)$ be the added kinase's share of the total weight. After the addition the entropy is $(1-q)H + h(q)$, where $h(q) = -q \ln q - (1-q)\ln(1-q)$ is the binary entropy, so the entropy does not decrease if and only if $H \le h(q)/q$. Because h is concave with $h(0) = 0$, $h(q)/q$ decreases in q. Since $q < r/(n + r)$ with $r = a/m \le e$, this gives $h(q)/q > \ln \frac{n+r}{r} + \frac{n}{r} \ln \frac{n+r}{n}$. As $H \le \ln n$, it suffices that $\ln n$ is at most this bound, which rearranges to $\ln r \le (1 + \frac{n}{r}) \ln(1 + \frac{r}{n})$. The right-hand side is at least 1 for all $n, r > 0$, because $\ln(1 + t) \ge t/(1 + t)$, and $\ln r \le 1$. The entropy therefore does not decrease, so the score does not increase. □ With $T = \tau^* - x_0$ the same argument covers every $\beta \in B$: the lightest gated kinase then has weight above $m = T \ln(1 + e^{(x_0 - \beta)/T})$, and $a/m < 2.213 < e$.

Sharpness. The bound cannot be lowered. With $\beta = x_0$ and $T < (\tau^* - x_0)/\ln(2^e - 1)$, a kinase added at τ^* has weight $r\,m$ with $r = \ln(1 + e^{(\tau^* - x_0)/T})/\ln 2 > e$, in units of the limiting weight $m = T \ln 2$ just above the floor. Take one active target with weight $s > r$ and n kinases just above the floor, with weights close to 1. For the limiting weights the entropy is $H_n = \ln(n + s) - s \ln s/(n + s)$, and after the addition $H'_n = \ln(n + s + r) - (s \ln s + r \ln r)/(n + s + r)$. Then $n(H'_n - H_n) \to r(1 - \ln r) < 0$ as $n \to \infty$, so for large enough n the entropy falls and the score rises; weights close to, but above, 1 keep this true with the gate strict. Numerically, profiles of this kind violate D3 by 3.1×10^{-5} bits on Klaeger at $T = 0.564$, and at $T = 1$ on Metz and $T = 10$ on Anastassiadis by up to 0.0014 and 0.014 bits. A random search over 20,000 profiles per dataset found no violation with $T = \tau^* - x_0$, and adding one kinase at τ^* never raises the score of a real compound. (`d4_entropy_vs_candidate.py`)

Validation. Table S5 reports the gated entropy's orientation (its correlation with the number of active kinases), D2 for its emphasis baseline and its gate, D2 for the entropy's baseline, and D3, the percentage of compounds whose score does not rise when one kinase is added at τ^*. Removing the non-binder gate inverts the orientation (gated $r = -0.946$, ungated $r = +0.91$ on Klaeger). (`candidate_validation.py`)

Table S5: The gated entropy: orientation, D2 (pairs reversing in B, %), and D3, over the n compounds that D1 allows to be scored.

Dataset	n	orientation r	D2, gate fixed	D2, gate swept	D2, entropy	D3 (%)
Davis	68	−0.968	0.44	9.22	7.90	100.0
Klaeger	200	−0.946	0.53	12.99	12.37	100.0
Anastassiadis	133	−0.934	0.36	14.39	10.58	100.0
Metz	699	−0.783	0.72	23.59	22.80	100.0

Variants. D3 is proved only for the Shannon order, either with $\beta = x_0$ and $T \ge (\tau^* - x_0)/\ln(2^e - 1) \approx 0.5816\,(\tau^* - x_0)$, or with any β in B and $T = \tau^* - x_0$. For other variants it is checked by adding

one kinase at τ^* to every scored compound, and it held for all of them. *Quantile floor.* Replacing the fixed floor with a per-compound low quantile of the profile (the 5th to 15th percentile) gives a variant invariant to shifts of the readout, but on these data it is not a separate test. On Davis every such percentile equals the floor for all but at most two compounds, and on Klaeger for every compound. At the 10th percentile the variant ranks compounds as the gated entropy does (Spearman 1.000), and across the five percentiles 0.04% (Davis) and 0% (Klaeger) of pairs reverse. *Rényi order.* Replacing the Shannon entropy by the Rényi entropy of order α, $-H_\alpha(p) = \frac{1}{\alpha-1}\log_2 \sum_i p_i^\alpha$, leaves D1 unchanged. Swept over B, the emphasis baseline reverses at most ε of pairs for orders 0.5 and 1 on both datasets and for $\alpha = 2$ on Davis, and more for $\alpha = 2$ on Klaeger and for $\alpha = 3$ (Table S6). Without the D1 gate (Section S5), every order from 0.01 to 3 reaches $p^* = 80$ on Klaeger and Metz, and the mean Spearman curves differ by at most 0.002 between orders. *Smoothing width.* On Klaeger, over $T = 0.5$, 0.75, 1, 1.5 and 2, orientation is unchanged ($r \approx -0.95$) and the rankings agree closely (worst-case pairwise Spearman 0.999). The emphasis baseline reverses 0.86%, 0.53%, 0.11% and 0.04% of pairs for $T = 0.75$, 1, 1.5 and 2, and 1.94% at $T = 0.5$, which is below the bound of the proposition; D3 held for every scored compound with one kinase added at τ^* at every width. As $T \to 0$ the softplus approaches the hard cutoff, giving back the entropy. (`family_construction.py`, `family_panelsize.py`, `candidate_validation.py`)

Table S6: Rényi-order variants on Davis and Klaeger: D2 for the emphasis baseline (pairs reversing in B, %), agreement with the Shannon ranking (Spearman), and p^* on Klaeger without the D1 gate (Section S5). (`family_construction.py`, `family_panelsize.py`)

Order	D2, Davis	D2, Klaeger	Spearman vs. $\alpha = 1$	p^*, Klaeger
$\alpha = 0.5$	0.40	0.14	0.999–1.000	80
$\alpha = 1$	0.44	0.53	—	80
$\alpha = 2$	0.92	1.04	0.999–1.000	80
$\alpha = 3$	1.14	1.65	0.998–0.999	80

Noise. With random noise added to the measured values, as described in the main text (Methods), the number of compounds whose D1 status changed in a draw was at most 4 (Klaeger, noise 0.1) and 11 (Klaeger, 0.3), at most 10 on Anastassiadis (9 percentage points) and 5 on Metz (0.3), and none on Davis. Over the compounds scored with and without noise, the gated entropy's ranking stayed at Spearman 0.978 or higher with the noise-free ranking (lowest on Anastassiadis), and the entropy's at 0.983 or higher. (`claim_checks.py` section 23)

S5: Recovery from subpanels

Without the D1 gate. Karaman et al. scored each compound's S-score on random subpanels and found that it varies more as panels shrink. Repeating that design for the four definitions, scoring every compound on every subpanel, mixes the ranking with each definition's convention for compounds that have no active target on the subpanel. For each panel size, in steps of 30 from 50 targets up to the full panel, we drew 50 random subsets, computed each definition on the subpanel with the dataset's floor x_0 as its baseline (entropy, Gini) or level (ratio, $k = 1$) and the activity threshold as the S-score cutoff, and correlated its ranking with the full-panel ranking. The S-score recovers the full-panel ranking fastest on every dataset, and the entropy as fast on Davis, Klaeger and Anastassiadis (Figure S4; p^* of 50 targets on Davis and Anastassiadis and 80 on Klaeger; on Metz 50 for the S-score and 80 for the entropy). The ratio is the slowest, needing 350 of 363 targets on Davis, 260 of 291 on Anastassiadis, 170 of 172 on Metz, and more than the largest subpanel tested on Klaeger ($r = 0.887$ at 200 targets). The Gini needs 110 targets on Davis and 200 on Anastassiadis and Klaeger, but matches the entropy on Metz (80). Its slower recovery comes from compounds whose Gini is undefined on the subpanel, which the convention ranks least selective. Excluding them gives the Gini p^* values of 50, 80, 50 and 80 targets. Changing the bar within $[0.80, 0.95]$ changes the p^* values but never reverses the order of two definitions on any dataset (Table S7). (`panel_size_analysis.py`, `pstar_bar_sensitivity.py`, `claim_checks.py` section 14)

Table S7: p^* (targets) without the D1 gate at four correlation bars. (`pstar_bar_sensitivity.py`)

Davis (363 targets)	entropy	Gini	S-score	ratio
bar = 0.80	50	80	50	320
bar = 0.85	50	110	50	320
bar = 0.90	50	110	50	350
bar = 0.95	50	170	80	350

Klaeger (217 targets)	entropy	Gini	S-score	ratio
bar = 0.80	50	170	50	200
bar = 0.85	50	200	80	200
bar = 0.90	80	200	80	> 200
bar = 0.95	110	200	140	> 200

Anastassiadis (291 targets)	entropy	Gini	S-score	ratio
bar = 0.80	50	170	50	260
bar = 0.85	50	170	50	260
bar = 0.90	50	200	50	260
bar = 0.95	80	260	80	290

Metz (172 targets)	entropy	Gini	S-score	ratio
bar = 0.80	50	50	50	170
bar = 0.85	50	50	50	170
bar = 0.90	80	80	50	170
bar = 0.95	110	110	80	170

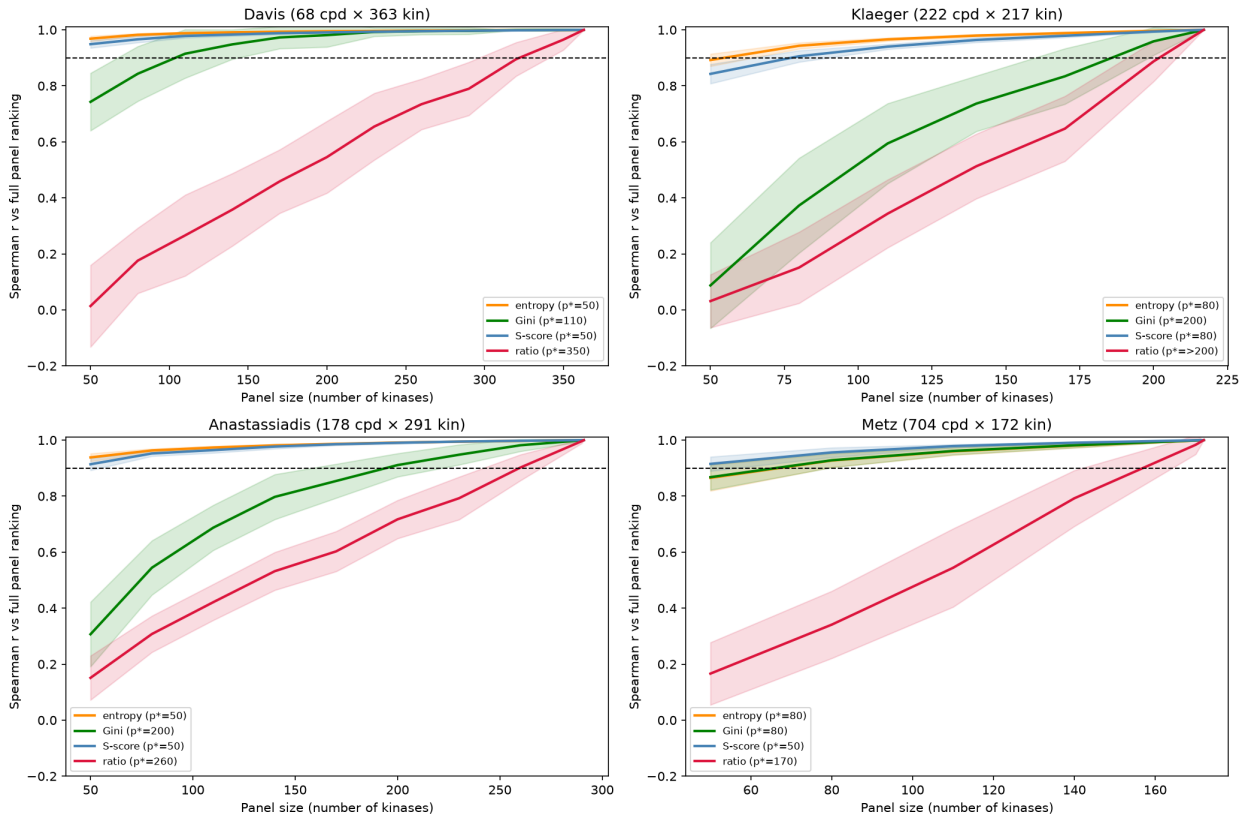


Figure S4: Mean Spearman correlation between rankings computed on random target subpanels and rankings on the full panel, as a function of subpanel size, with every compound scored on every subpanel. Shaded bands: ±1 standard deviation across 50 subsamples. Dashed line: $r = 0.90$; p^* (legend) is the smallest panel reaching mean $r > 0.90$.

Variability across subpanels. In the main-text recovery analysis, with D1 applied, p^* is a threshold on the mean over 50 subpanels. At p^*, the share of individual subpanels above 0.90 is 70%, 72%, 86% and 62% for the gated entropy on Davis, Klaeger, Anastassiadis and Metz, 64–90% for the S-score at x_0, 66–82% for the entropy at x_0, 62–74% for the Gini, 60–80% for the S-score at τ^*, and 78–90% for the entropy at τ^*. Keeping floor-coded pairs at the floor instead of giving them zero weight leaves the published entropy's p^* unchanged (300, 200 and 150 targets). With bars of 0.80, 0.85 or 0.95 instead of 0.90, the share-following scores (the S-score and entropy at x_0, the Gini and the gated entropy) reach the bar no later than the S-score and the entropy at τ^* on Davis, Klaeger and Anastassiadis, and the ratio, the published entropy and $P_{\max}$ are the slowest scores on every dataset. On a 50-target subpanel, a compound with one active target stays scored with probability 0.14 (Davis), 0.23 (Klaeger), 0.17 (Anastassiadis) and 0.29 (Metz), and one with ten active targets with probability 0.78, 0.93, 0.85 and 0.97. (`panel_recovery.py`)

S6: Cross-dataset and supplementary comparisons

Dataset and definition. On the 46 compounds profiled in both Davis and Klaeger and the 201 kinases both panels share, the same definition's Davis and Klaeger rankings correlate at 0.78 (S-score), 0.84 (entropy), 0.45 (Gini) and 0.34 (ratio, $p = 0.02$), and different definitions within one dataset at −0.08–1.00 (mean 0.57). Kinases are matched by gene symbol, with renamed Davis symbols mapped to current ones (for example, Davis's MLCK, accession NP_872299.2, is MYLK3); a check of the other renamed symbols found no further matches. Klaeger has six compounds with no active target on this shared panel. Among the 40 compounds active in both, the same definition's rankings correlate at 0.74, 0.85, 0.84 and 0.46 (mean 0.72, against 0.73 between definitions). On the full panels the cross-dataset correlations are 0.14–0.81 (mean 0.53). A compound's rank instability within a dataset is not associated with how far its rank moves between the datasets ($|r| \leq 0.18$). (`davis_klaeger_overlap.py`, `platform_vs_definition.py`, `claim_checks.py` section 9)

Bosc et al.'s scores and a CATDS proxy. Bosc et al.'s window score (WS) and ranking score (RS) correlate strongly with the ratio on every dataset (Table S8). RS at rank $k + 1$ is $R(k)$, so its agreement with the ratio follows from its definition. The CATDS score scores one designated target at a chosen concentration from Kinobeads curves, which only Klaeger provides. A proxy, the top target's share of the baseline-subtracted profile, correlates more with the Gini than with the ratio on every dataset, and more with the entropy than with the ratio except on Klaeger. It lacks the concentration dependence of the published score. (`bosc_catds_comparison.py`)

Table S8: Spearman correlation of WS and RS with the ratio, and of the CATDS proxy with the ratio, entropy and Gini.

Dataset	n_{cpd}	WS vs. ratio	RS vs. ratio	CATDS vs. ratio	CATDS vs. entropy / Gini
Davis	68	0.932	0.904	0.636	0.977 / 0.977
Klaeger	222	0.926	0.890	0.724	0.637 / 0.982
Anastassiadis	178	0.957	0.883	0.698	0.945 / 0.964
Metz	704	0.910	0.899	0.447	0.883 / 0.906

S7: Clinical analysis

For the FDA-approved Klaeger compounds, we correlated each definition's median selectivity rank with the proportions of FAERS reports that were serious, fatal or involved hospitalization ($n = 46$) and with discontinuation rates reported in FDA labels ($n = 33$). FAERS counts were pulled from the openFDA adverse-event interface (`faers_pull.py`); discontinuation rates were extracted by hand from label text (`clinical_safety_data.csv`). No correlation was significant (all $|r| \leq 0.20$, all $p > 0.19$). Raw report counts correlate weakly with the ratio ranking ($r = 0.30$, $p = 0.04$), which may reflect prescription volume. Grade 3/4 adverse-event rates were available for too few drugs. FAERS proportions describe reports, not incidence among treated patients; label discontinuation rates come from trials with different populations and doses; and both depend on indication and prescription volume. The analysis therefore cannot isolate an effect of selectivity on toxicity. (`clinical_outcomes_correlation.py`)

S8: Reproducibility index

Every result in the main text is printed by a script in the code repository. Table S9 gives, for each result, the script that computes it. Each script saves its printed output as `outputs/<script>.txt`; numbers in parentheses are numbered sections of that output. Random steps are seeded, so rerunning a script reproduces its output.

Table S9: Provenance of reported results.

Location	Result	Script (output section)
Methods; Table 1; S1	Curated panels, curation tables, Klaeger values below the floor	`curate_panels.py`
Methods; S1	Active fractions, Anastassiadis weak-inhibition cutoff, Metz coverage before restriction, Klaeger compounds with few active kinases, Davis stricter threshold	`dataset_descriptive_stats.py`
Methods; S1	Metz coverage of the retained compounds	`metz_coverage.py`, `claim_checks.py` (10)
Methods; S2	Mean instead of median ranks	`claim_checks.py` (3)
Table 3; Results, Agreement	Correlations between definitions; Klaeger active compounds; bootstrap	`selectivity_analysis.py`, `klaeger_analysis.py`, `claim_checks.py` (2)
Table 3; Figure 1; Table S2	Published entropy and Gini, both floor treatments; every measured Klaeger affinity retained	`original_operationalization_robustness.py`, `metric_fidelity_robustness.py`
Results, Agreement; S3	Compounds with no active kinase: ranks, aggregated and per-definition instability, undefined configurations	`claim_checks.py` (13, 20), `failure_mode_illustration.py`, `undefined_score_convention.py`
Results, Agreement; S3	Targets of the compounds with no detected binding	`claim_checks.py` (1)
Results, Agreement; Figure S1; S2	Four datasets, all and active compounds; Metz coverage dependence	`additional_datasets_analysis.py`, `claim_checks.py` (10, 11)
Methods; Results, Agreement; Table S1	Released against curated panels	`davis_gene_level.py`, `klaeger_protein_kinases.py`
Results, Agreement; Discussion; S6	Dataset against definition on shared compounds and kinases (renamed symbols mapped); active fractions; instability against cross-dataset rank change	`davis_klaeger_overlap.py`, `platform_vs_definition.py`, `claim_checks.py` (9)
Results, Agreement; Results, Recovery; S1	Klaeger confidence filter, including gated p^*	`klaeger_filter_sensitivity.py`
Results, Sensitivity; Table 4; Table S3; Figure S2	Pairs reversing in B (D2 verdicts); gate against baseline	`d3_pair_reversals.py`, `failure_mode_illustration.py`
Results, Sensitivity	Ratio at near-tied top targets	`d2_near_tie_sensitivity.py`
Figure 2	Binding profiles of two Davis compounds	`binding_profiles.py`

Location	Result	Script (output section)
Results, Sensitivity; Table S4; Figure S3	Profile features and instability, both ranking populations; compounds defined at every baseline; Anastassiadis and Metz relations	`instability_by_family.py`, `additional_datasets_analysis.py`, `claim_checks.py` (4, 8, 15–17, 21)
S3	Undefined features; further features	`feature_placeholder_sensitivity.py`, `instability_predictor_candidates.py`
Results, Sensitivity; S3	Broad profiles; most unstable and most stable compounds	`claim_checks.py` (5–7, 18)
Methods; Results, Criteria; S4	Bound on T; entropy and gated entropy under weak addition; worked example; discontinuity; strengthening example; sharpness	`d4_entropy_vs_candidate.py`
Results, Criteria	Realistic panel expansion	`claim_checks.py` (19)
Results, Criteria; Table S5	Orientation, D3 at τ^*, smoothing width	`candidate_validation.py`
Results, Criteria	Agreement with the entropy	`candidate_measure.py`
Results, Criteria; S4	Dependence on τ^*; noise	`claim_checks.py` (22, 23)
Results, Criteria; Results, Recovery; Figure 3; S5	Coverage, recoverability, conditional agreement, variability, influence, agreement with the S-score at x_0; other bars; probability of staying scored	`panel_recovery.py`
S4; Table S6	Variants of the gated entropy	`family_construction.py`, `family_panelsize.py`
S5; Figure S4; Table S7	Recovery without the D1 gate	`panel_size_analysis.py`, `pstar_bar_sensitivity.py`, `claim_checks.py` (14)
Discussion; S6; Table S8	Bosc et al.'s scores and the CATDS proxy	`bosc_catds_comparison.py`
Discussion; S7	Clinical outcomes	`clinical_outcomes_correlation.py`, `faers_pull.py`